

WORKING WITH RELATIONAL DATABASES EXERCISE MANUAL



Fidelity LEAP
Technology Immersion Program

This page was intentionally left blank.

Table of Contents

| | |
|---|-----------|
| Chapter 1: What Is Structured Query Language? | 1 |
| Exercise 1.1: Using SQL Developer | 1 |
| Chapter 2: SQL Query Syntax | 3 |
| Exercise 2.1: Selecting Data | 3 |
| Chapter 3: SQL Scalar Functions | 9 |
| Exercise 3.1: Using Scalar Functions | 9 |
| Chapter 4: SQL Joins | 13 |
| Exercise 4.1: Working with <code>INNER JOINS</code> | 13 |
| Exercise 4.2: Using <code>OUTER JOINS</code> | 17 |
| Chapter 5: Additional SQL Functions | 23 |
| Exercise 5.1: Additional SQL Functions | 23 |
| Chapter 6: Data Manipulation Language | 25 |
| Exercise 6.1: Manipulating Data | 25 |
| Chapter 7: Databases with JDBC (Java Database Connectivity) | 27 |
| Exercise 7.1: Connecting to a Database | 27 |
| Exercise 7.2: Creating Objects from a Database Query | 29 |
| Chapter 8: Updating Databases | 33 |
| Exercise 8.1: Inserting a Record with JDBC | 33 |
| Exercise 8.2: Updating and Deleting Records | 35 |
| Exercise 8.3: Working with Multi-Table Queries | 36 |
| Chapter 9: Working with a Data Access Object | 37 |
| Exercise 9.1: Debugging Data Access Object Methods | 37 |
| Optional Exercise 9.2: Testing a Business Service with a Mock DAO | 38 |
| Exercise 9.3 Putting It All Together | 40 |
| Chapter 10: Advanced JDBC | 41 |
| Optional Exercise 10.1: JDBC Code Review | 41 |
| Optional Exercise 10.2: One-to-Many Queries | 42 |
| Optional Exercise 10.3: Testing a Service Proxy with Mockito | 43 |
| Optional Exercise 10.4: Writing an Integration Test for the Service and DAO | 46 |
| Chapter 11: Aggregating Information | 49 |
| Exercise 11.1: Using the Aggregate Functions | 49 |
| Exercise 11.2: <code>GROUP BY</code> and <code>HAVING</code> | 51 |
| Exercise 11.3: Using Subqueries | 54 |
| Chapter 12: Set Operators | 57 |
| Exercise 12.1: Set Operators | 57 |

| | |
|--|-----------|
| Chapter 13: Programming with PL/SQL | 59 |
| Exercise 13.1: Building Anonymous Blocks | 59 |
| Exercise 13.2: Using Cursors | 61 |
| Chapter 14: Creating Stored Procedures, Functions, and Packages | 63 |
| Exercise 14.1: Stored Procedures, Functions, and Packages | 63 |
| Chapter 15: Testing PL/SQL | 65 |
| Exercise 15.1: Writing PL/SQL Tests with utPLSQL | 65 |
| Exercise 15.2: Testing Updates With utPLSQL | 67 |
| Chapter 16: Creating Triggers..... | 69 |
| Exercise 16.1: Working with Triggers | 69 |
| Chapter 17: Data Definition Language | 71 |
| Exercise 17.1: Table Management | 71 |
| Chapter 20: Amazon DynamoDB | 75 |
| Exercise 20.1: Access DynamoDB Using AWS CLI | 75 |
| Exercise 20.2: Java Document API | 78 |
| Exercise 20.3: Java Object Mapper | 79 |

Chapter 1: What Is Structured Query Language?

Exercise 1.1: Using SQL Developer

For this exercise, we will become familiar with SQL Developer.

1. Locate the SQL Developer executable on the desktop.
2. Double-click to see that it starts correctly.
3. Right-click the **HR** connection and select Properties from the menu. Examine the parameters. Where is the Oracle database server running?
4. Make the following configuration changes in the Tools | Preferences Menu:
 - a. In the Code Editor Line Gutter section, choose the option to **Show Line Numbers**.

Tools / Preferences / Code Editor / Line Gutter

- b. In the Database NLS Parameters section, modify the Date Format to display **DD-MON-YYYY**.

Tools / Preferences / Database / NLS Parameters

5. Open the **HR** connection.

You can do this either by clicking the + (plus sign) to the left of the Connection Name or by right-clicking the Connection icon and selecting Connect.

6. Expand the **Tables** branch to show all of the tables from **HR**.

Click the + (plus sign) or double-click the name.

7. How many tables are owned by **HR**?

-
8. Open SQL Worksheet for **HR**.

Use the SQL Worksheet icon in the toolbar.

9. Enter the following statement into the SQL Worksheet window and execute it:

```
SELECT * FROM locations;
```

Click the Run Statement icon or the <F9> key to run the statement.

10. Execute the same statement as a script and note the difference.

Click the Run Script icon or the <F5> key to run the script.

11. Save the SQL script for later use as a file called `loc.sql`.
12. Enter the following statement underneath the existing statement.

```
SELECT * FROM countries;
```

13. Execute both statements by clicking the <F9> key and note the result.
14. Now, execute both statements together as a script (<F5>). Note that the results of both statements are displayed in the Script Output.
15. Clear the contents of the SQL Worksheet.
16. Open the `loc.sql` file.

Click the Open File icon or use the File menu and select Open.

17. Run the SQL statement.

Congratulations! You have completed this exercise.

Chapter 2: SQL Query Syntax

Exercise 2.1: Selecting Data

Connect to the SCOTT account

Unless the order is specified, the order of your results may differ.

1. Write a query to display the dname and deptno of all rows in the dept table.

| DNAME | DEPTNO |
|------------|--------|
| ACCOUNTING | 10 |
| RESEARCH | 20 |
| SALES | 30 |
| OPERATIONS | 40 |

2. Write a query to display ALL of the columns and rows in the dept table.

| DEPTNO | DNAME | LOC |
|--------|------------|----------|
| 10 | ACCOUNTING | NEW YORK |
| 20 | RESEARCH | DALLAS |
| 30 | SALES | CHICAGO |
| 40 | OPERATIONS | BOSTON |

3. Write a query to display the dname, deptno, and location of all rows in the dept table labeling them Name, DEPT# and Dept Location, respectively.

| Name | DEPT# | Dept Location |
|------------|-------|---------------|
| ACCOUNTING | 10 | NEW YORK |
| RESEARCH | 20 | DALLAS |
| SALES | 30 | CHICAGO |
| OPERATIONS | 40 | BOSTON |

4. Write a query to display the deptno of each row in the emp table.

| DEPTNO |
|--------|
| 20 |
| 30 |
| 30 |
| 20 |
| 30 |
| 30 |
| 10 |
| 20 |
| 10 |
| 30 |
| 20 |
| 30 |
| 20 |
| 30 |
| 20 |
| 10 |

Exercise Manual

5. Write a query to display each deptno in the emp table only once.

```
DEPTNO
-----
      30
      10
      20
```

6. Write a query to display the deptno and job of each row in the emp table.

```
DEPTNO  JOB
-----  -
      20  CLERK
      30  SALESMAN
      30  SALESMAN
      20  MANAGER
      30  SALESMAN
      30  MANAGER
      10  MANAGER
      20  ANALYST
      10  PRESIDENT
      30  SALESMAN
      20  CLERK
      30  CLERK
      20  ANALYST
      10  CLERK
```

7. Write a query to display each unique combination of deptno and job in the emp table.

```
DEPTNO  JOB
-----  -
      20  MANAGER
      20  ANALYST
      10  PRESIDENT
      10  CLERK
      30  SALESMAN
      10  MANAGER
      20  CLERK
      30  MANAGER
      30  CLERK
```

8. Write a query to display the names of the employees who work for dept 30 in the emp table.

```
ENAME
-----
ALLEN
WARD
MARTIN
BLAKE
TURNER
JAMES
```


Exercise Manual

9. Write a query to display the names of the employees who were hired on Dec. 17, 1981. Specify the date in a safe format.

```
no rows selected
```

10. Write a query to display the names of the employees who were hired on or after Dec. 17, 1981.

```
ENAME
-----
SCOTT
ADAMS
MILLER
```

11. Write a query to display the names of the employees who have the job of clerk.

```
no rows selected
```

12. Write a query to display the names of the employees who have the job of CLERK.

```
ENAME
-----
SMITH
ADAMS
JAMES
MILLER
```

13. Write a query to display the names of the employees whose salary is greater than 2500.

```
ENAME
-----
JONES
BLAKE
SCOTT
KING
FORD
```

14. Write a query to display the names of the employees whose salary is in the range (inclusive) of 1000 and 1600.

```
NAME
-----
ALLEN
WARD
MARTIN
TURNER
ADAMS
MILLER
```

Exercise Manual

15. Write a query to display the names of the employees whose names contain "ER".

```
ENAME
-----
TURNER
MILLER
```

16. Write a query to display the names and employee numbers of the employees whose commission is undefined.

```
EMPNO ENAME
-----
7369 SMITH
7566 JONES
7698 BLAKE
7782 CLARK
7788 SCOTT
7839 KING
7876 ADAMS
7900 JAMES
7902 FORD
7934 MILLER
```

17. Write a query to display the names, employee numbers, and commissions of the employees, sequencing the data in commission ascending order.

| | EMPNO | ENAME | COMM |
|----|-------|--------|--------|
| 1 | 7844 | TURNER | 0 |
| 2 | 7499 | ALLEN | 300 |
| 3 | 7521 | WARD | 500 |
| 4 | 7654 | MARTIN | 1400 |
| 5 | 7788 | SCOTT | (null) |
| 6 | 7839 | KING | (null) |
| 7 | 7876 | ADAMS | (null) |
| 8 | 7900 | JAMES | (null) |
| 9 | 7902 | FORD | (null) |
| 10 | 7934 | MILLER | (null) |
| 11 | 7698 | BLAKE | (null) |
| 12 | 7566 | JONES | (null) |
| 13 | 7369 | SMITH | (null) |
| 14 | 7782 | CLARK | (null) |

Exercise Manual

18. Write a query to display the names, employee numbers, and commissions of the employees sequencing the data in commission descending order.

| | E... | ENAME | COMM |
|----|------|--------|--------|
| 1 | 7369 | SMITH | (null) |
| 2 | 7782 | CLARK | (null) |
| 3 | 7902 | FORD | (null) |
| 4 | 7900 | JAMES | (null) |
| 5 | 7876 | ADAMS | (null) |
| 6 | 7566 | JONES | (null) |
| 7 | 7698 | BLAKE | (null) |
| 8 | 7934 | MILLER | (null) |
| 9 | 7788 | SCOTT | (null) |
| 10 | 7839 | KING | (null) |
| 11 | 7654 | MARTIN | 1400 |
| 12 | 7521 | WARD | 500 |
| 13 | 7499 | ALLEN | 300 |
| 14 | 7844 | TURNER | 0 |

19. Write a query to display the names and employee numbers of the employees sequencing the data in commission descending order, forcing those with unknown commissions to the bottom of the list.

| | EMPNO | ENAME | COMM |
|----|-------|--------|--------|
| 1 | 7654 | MARTIN | 1400 |
| 2 | 7521 | WARD | 500 |
| 3 | 7499 | ALLEN | 300 |
| 4 | 7844 | TURNER | 0 |
| 5 | 7788 | SCOTT | (null) |
| 6 | 7839 | KING | (null) |
| 7 | 7876 | ADAMS | (null) |
| 8 | 7900 | JAMES | (null) |
| 9 | 7902 | FORD | (null) |
| 10 | 7934 | MILLER | (null) |
| 11 | 7698 | BLAKE | (null) |
| 12 | 7566 | JONES | (null) |
| 13 | 7369 | SMITH | (null) |
| 14 | 7782 | CLARK | (null) |

Congratulations! You have completed this exercise.

This page intentionally left blank.

Chapter 3: SQL Scalar Functions

Exercise 3.1: Using Scalar Functions

Unless the order is specified, the order of your results may differ.

Connect to the HR account.

1. Write a query to display the first name, last name, and salary of all employees in department 30, formatting the salary with commas and a floating dollar sign.

| FIRST_NAME | LAST_NAME | SALARY |
|------------|------------|----------|
| Den | Raphaely | \$11,000 |
| Alexander | Khoo | \$3,100 |
| Shelli | Baida | \$2,900 |
| Sigal | Tobias | \$2,800 |
| Guy | Himuro | \$2,600 |
| Karen | Colmenares | \$2,500 |

2. Write a query to display the first name, last name, and date hired of all employees in department 30, formatting the date to be year-month#-day.

| FIRST_NAME | LAST_NAME | Date Hired |
|------------|------------|------------|
| Den | Raphaely | 2002-12-07 |
| Alexander | Khoo | 2003-05-18 |
| Shelli | Baida | 2005-12-24 |
| Sigal | Tobias | 2005-07-24 |
| Guy | Himuro | 2006-11-15 |
| Karen | Colmenares | 2007-08-10 |

3. Write a query to display the salary of all employees in department 30. Also show the salary rounded and truncated to thousands.

| FIRST_NAME | LAST_NAME | RDSAL | TSAL | SALARY |
|------------|------------|-------|-------|--------|
| Den | Raphaely | 11000 | 11000 | 11000 |
| Alexander | Khoo | 3000 | 3000 | 3100 |
| Shelli | Baida | 3000 | 2000 | 2900 |
| Sigal | Tobias | 3000 | 2000 | 2800 |
| Guy | Himuro | 3000 | 2000 | 2600 |
| Karen | Colmenares | 3000 | 2000 | 2500 |

Exercise Manual

- Write a query to display names of all employees in department 30. Their first name should be in lower case; their last name in upper case. Sequence the list in (ascending) first name, last name order.

| LNAME | UNAME |
|-----------|------------|
| alexander | KHOO |
| den | RAPHAELY |
| guy | HIMURO |
| karen | COLMENARES |
| shelll | BAIDA |
| sigal | TOBIAS |

- Write a query to display the initial of the first name followed by a period followed by the last name of all employees in department 30. Sequence the list in alphabetical order of this formatted name.

| NAME |
|---------------|
| A. Khoo |
| D. Raphaely |
| G. Himuro |
| K. Colmenares |
| S. Baida |
| S. Tobias |

- Write a query to display the street address, followed by the street address stripped of any leading numeric digits, spaces, or dashes (Street Name) for all rows in the locations table. Order the list by the Street Name.

| STREET_ADDRESS | Street Name |
|--|--|
| 8204 Arthur St | Arthur St |
| 6092 Boxwood St | Boxwood St |
| 93091 Calle della Testa | Calle della Testa |
| 2004 Charade Rd | Charade Rd |
| 9702 Chester Road | Chester Road |
| 198 Clementi North | Clementi North |
| 2011 Interiors Blvd | Interiors Blvd |
| 2014 Jabberwocky Rd | Jabberwocky Rd |
| 9450 Kamiya-cho | Kamiya-cho |
| 40-5-12 Laogianggen | Laogianggen |
| Magdalen Centre, The Oxford Science Park | Magdalen Centre, The Oxford Science Park |
| Mariano Escobedo 9991 | Mariano Escobedo 9991 |
| Murtenstrasse 921 | Murtenstrasse 921 |
| Pieter Breughelstraat 837 | Pieter Breughelstraat 837 |
| Rua Frei Caneca 1360 | Rua Frei Caneca 1360 |
| 20 Rue des Corps-Saints | Rue des Corps-Saints |
| Schwanthalerstr. 7031 | Schwanthalerstr. 7031 |
| 2017 Shinjuku-ku | Shinjuku-ku |
| 147 Spadina Ave | Spadina Ave |
| 1297 Via Cola di Rie | Via Cola di Rie |
| 12-98 Victoria Street | Victoria Street |
| 1298 Vileparle (E) | Vileparle (E) |
| 2007 Zagora St | Zagora St |

Exercise Manual

- Write a query to display the street address, followed by length of the street address (Street Length) for all rows in the `locations` table. Sequence the list in the Street Length order.

| STREET_ADDRESS | Street Length |
|--|---------------|
| 8204 Arthur St | 14 |
| 2007 Zagora St | 14 |
| 2004 Charade Rd | 15 |
| 9450 Kamiya-cho | 15 |
| 6092 Boxwood St | 15 |
| 147 Spadina Ave | 15 |
| 2017 Shinjuku-ku | 16 |
| Murtenstrasse 921 | 17 |
| 9702 Chester Road | 17 |
| 1298 Vileparle (E) | 18 |
| 198 Clementi North | 18 |
| 2014 Jabberwocky Rd | 19 |
| 40-5-12 Laogianggen | 19 |
| 2011 Interiors Blvd | 19 |
| 1297 Via Cola di Rie | 20 |
| Schwanthalerstr. 7031 | 21 |
| 12-98 Victoria Street | 21 |
| Rua Frei Caneca 1360 | 21 |
| Mariano Escobedo 9991 | 21 |
| 20 Rue des Corps-Saints | 23 |
| 93091 Calle della Testa | 23 |
| Pieter Breughelstraat 837 | 25 |
| Magdalen Centre, The Oxford Science Park | 40 |

- Write a query to display the location ID, the street address, city, and state province of all rows in the `locations` table that contain either the string "RUA" or "RUE" in the street address. Sequence the list in descending sequence on location ID.

| LOCATION_ID | STREET_ADDRESS | CITY | STATE_PROVINCE |
|-------------|-------------------------|-----------|----------------|
| 2900 | 20 Rue des Corps-Saints | Geneva | Geneve |
| 2800 | Rua Frei Caneca 1360 | Sao Paulo | Sao Paulo |

Congratulations! You have completed this exercise.

This page intentionally left blank.

Chapter 4: SQL Joins

Exercise 4.1: Working with INNER JOINS

Save your queries for later use.

Connect to the HR account.

1. Joining the `locations` and `departments` tables, display the city, location ID, and department name.

| CITY | LOCATION_ID | DEPARTMENT_NAME |
|---------------------|-------------|----------------------|
| Southlake | 1400 | IT |
| South San Francisco | 1500 | Shipping |
| Seattle | 1700 | Administration |
| Seattle | 1700 | Purchasing |
| Seattle | 1700 | Executive |
| Seattle | 1700 | Finance |
| Seattle | 1700 | Accounting |
| Seattle | 1700 | Treasury |
| Seattle | 1700 | Corporate Tax |
| Seattle | 1700 | Control And Credit |
| Seattle | 1700 | Shareholder Services |
| Seattle | 1700 | Benefits |
| Seattle | 1700 | Manufacturing |
| Seattle | 1700 | Construction |
| Seattle | 1700 | Contracting |
| Seattle | 1700 | Operations |
| Seattle | 1700 | IT Support |
| Seattle | 1700 | NOC |
| Seattle | 1700 | IT Helpdesk |
| Seattle | 1700 | Government Sales |
| Seattle | 1700 | Retail Sales |
| Seattle | 1700 | Recruiting |
| Seattle | 1700 | Payroll |
| Toronto | 1800 | Marketing |
| London | 2400 | Human Resources |
| Oxford | 2500 | Sales |
| Munich | 2700 | Public Relations |

Exercise Manual

2. Joining the `locations` and `countries` tables, display the country name and city.

| COUNTRY_NAME | CITY |
|--------------------------|---------------------|
| Australia | Sydney |
| Brazil | Sao Paulo |
| Canada | Toronto |
| Canada | Whitehorse |
| Switzerland | Geneva |
| Switzerland | Bern |
| China | Beijing |
| Germany | Munich |
| India | Bombay |
| Italy | Roma |
| Italy | Venice |
| Japan | Tokyo |
| Japan | Hiroshima |
| Mexico | Mexico City |
| Netherlands | Utrecht |
| Singapore | Singapore |
| United Kingdom | London |
| United Kingdom | Oxford |
| United Kingdom | Stretford |
| United States of America | Southlake |
| United States of America | South San Francisco |
| United States of America | South Brunswick |
| United States of America | Seattle |

Exercise Manual

3. Joining the `locations`, `countries`, and `departments` tables, display the country name, city, and department name.

| COUNTRY_NAME | CITY | DEPARTMENT_NAME |
|--------------------------|---------------------|----------------------|
| United States of America | Southlake | IT |
| United States of America | South San Francisco | Shipping |
| United States of America | Seattle | Administration |
| United States of America | Seattle | Purchasing |
| United States of America | Seattle | Executive |
| United States of America | Seattle | Finance |
| United States of America | Seattle | Accounting |
| United States of America | Seattle | Treasury |
| United States of America | Seattle | Corporate Tax |
| United States of America | Seattle | Control And Credit |
| United States of America | Seattle | Shareholder Services |
| United States of America | Seattle | Benefits |
| United States of America | Seattle | Manufacturing |
| United States of America | Seattle | Construction |
| United States of America | Seattle | Contracting |
| United States of America | Seattle | Operations |
| United States of America | Seattle | IT Support |
| United States of America | Seattle | NOC |
| United States of America | Seattle | IT Helpdesk |
| United States of America | Seattle | Government Sales |
| United States of America | Seattle | Retail Sales |
| United States of America | Seattle | Recruiting |
| United States of America | Seattle | Payroll |
| Canada | Toronto | Marketing |
| United Kingdom | London | Human Resources |
| United Kingdom | Oxford | Sales |
| Germany | Munich | Public Relations |

4. Joining the `employees` and `job_history` tables, display the employee ID, first and last name, and the job ID. Display the output in sequence by `employee_id`.

| EMPLOYEE_ID | FIRST_NAME | LAST_NAME | JOB_ID |
|-------------|------------|-----------|------------|
| 101 | Neena | Kochhar | AC_MGR |
| 101 | Neena | Kochhar | AC_ACCOUNT |
| 102 | Lex | De Haan | IT_PROG |
| 114 | Den | Raphaely | ST_CLERK |
| 122 | Payam | Kaufling | ST_CLERK |
| 176 | Jonathon | Taylor | SA_REP |
| 176 | Jonathon | Taylor | SA_MAN |
| 200 | Jennifer | Whalen | AD_ASST |
| 200 | Jennifer | Whalen | AC_ACCOUNT |
| 201 | Michael | Hartstein | MK_REP |

Exercise Manual

5. Joining the `jobs` and `job_history` tables, display the job title, employee ID, and starting date for all employees who started in that job after Jan. 1, 1998.

| JOB_TITLE | EMPLOYEE_ID | START_DATE |
|--------------------------|-------------|-------------|
| Public Accountant | 200 | 01-JUL-2002 |
| Accounting Manager | 101 | 28-OCT-2001 |
| Programmer | 102 | 13-JAN-2001 |
| Marketing Representative | 201 | 17-FEB-2004 |
| Sales Manager | 176 | 01-JAN-2007 |
| Sales Representative | 176 | 24-MAR-2006 |
| Stock Clerk | 114 | 24-MAR-2006 |
| Stock Clerk | 122 | 01-JAN-2007 |

6. Modify the above query: remove the start date restriction and also include the employees' first and last names.

| JOB_TITLE | EMPLOYEE_ID | START_DATE | FIRST_NAME | LAST_NAME |
|--------------------------|-------------|-------------|------------|-----------|
| Accounting Manager | 101 | 28-OCT-2001 | Neena | Kochhar |
| Public Accountant | 101 | 21-SEP-1997 | Neena | Kochhar |
| Programmer | 102 | 13-JAN-2001 | Lex | De Haan |
| Stock Clerk | 114 | 24-MAR-2006 | Den | Raphaely |
| Stock Clerk | 122 | 01-JAN-2007 | Payam | Kaufling |
| Sales Representative | 176 | 24-MAR-2006 | Jonathon | Taylor |
| Sales Manager | 176 | 01-JAN-2007 | Jonathon | Taylor |
| Administration Assistant | 200 | 17-SEP-1995 | Jennifer | Whalen |
| Public Accountant | 200 | 01-JUL-2002 | Jennifer | Whalen |
| Marketing Representative | 201 | 17-FEB-2004 | Michael | Hartstein |

Congratulations! You have completed this exercise.

Exercise 4.2: Using OUTER JOINS

Connect to the HR account.

Using the standard outer join syntax, write queries to display the following information:

Exercise Manual

1. Joining the `employees` and `job_history` tables, display the employee ID, first and last name, and the job ID. Include all employees, whether or not they have any job history. Display in employee ID order. Modify your solution to Step 4 of the previous exercise.

| | EMPLOYEE_ID | FIRST_NAME | LAST_NAME | JOB_ID |
|----|-------------|-------------|-------------|------------|
| 1 | 100 | Steven | King | (null) |
| 2 | 101 | Neena | Kochhar | AC_ACCOUNT |
| 3 | 101 | Neena | Kochhar | AC_MGR |
| 4 | 102 | Lex | De Haan | IT_PROG |
| 5 | 103 | Alexander | Hunold | (null) |
| 6 | 104 | Bruce | Ernst | (null) |
| 7 | 105 | David | Austin | (null) |
| 8 | 106 | Valli | Pataballa | (null) |
| 9 | 107 | Diana | Lorentz | (null) |
| 10 | 108 | Nancy | Greenberg | (null) |
| 11 | 109 | Daniel | Faviet | (null) |
| 12 | 110 | John | Chen | (null) |
| 13 | 111 | Ismael | Sciarra | (null) |
| 14 | 112 | Jose Manuel | Urman | (null) |
| 15 | 113 | Luis | Popp | (null) |
| 16 | 114 | Den | Raphaely | ST_CLERK |
| 17 | 115 | Alexander | Khoo | (null) |
| 18 | 116 | Shelli | Baida | (null) |
| 19 | 117 | Sigal | Tobias | (null) |
| 20 | 118 | Guy | Himuro | (null) |
| 21 | 119 | Karen | Colmenares | (null) |
| 22 | 120 | Matthew | Weiss | (null) |
| 23 | 121 | Adam | Fripp | (null) |
| 24 | 122 | Payam | Kaufling | ST_CLERK |
| 25 | 123 | Shanta | Vollman | (null) |
| 26 | 124 | Kevin | Mourgos | (null) |
| 27 | 125 | Julia | Nayer | (null) |
| 28 | 126 | Irene | Mikkilineni | (null) |
| 29 | 127 | James | Landry | (null) |
| 30 | 128 | Steven | Markle | (null) |
| 31 | 129 | Laura | Bissot | (null) |
| 32 | 130 | Mozhe | Atkinson | (null) |
| 33 | 131 | James | Marlow | (null) |
| 34 | 132 | TJ | Olson | (null) |
| 35 | 133 | Jason | Mallin | (null) |
| 36 | 134 | Michael | Rogers | (null) |
| 37 | 135 | Ki | Gee | (null) |
| 38 | 136 | Hazel | Philtanker | (null) |
| 39 | 137 | Renske | Ladwig | (null) |
| 40 | 138 | Stephen | Stiles | (null) |
| 41 | 139 | John | Seo | (null) |
| 42 | 140 | Joshua | Patel | (null) |
| 43 | 141 | Trenna | Rajs | (null) |
| 44 | 142 | Curtis | Davies | (null) |
| 45 | 143 | Randall | Matos | (null) |
| 46 | 144 | Peter | Vargas | (null) |
| 47 | 145 | John | Russell | (null) |
| 48 | 146 | Karen | Partners | (null) |
| 49 | 147 | Alberto | Errazuriz | (null) |

Continued on the next page

Exercise Manual

| | | | |
|----|-----------------|------------|--------|
| 50 | 148 Gerald | Cambrault | (null) |
| 51 | 149 Eleni | Zlotkey | (null) |
| 52 | 150 Peter | Tucker | (null) |
| 53 | 151 David | Bernstein | (null) |
| 54 | 152 Peter | Hall | (null) |
| 55 | 153 Christopher | Olsen | (null) |
| 56 | 154 Nanette | Cambrault | (null) |
| 57 | 155 Oliver | Tuvault | (null) |
| 58 | 156 Janette | King | (null) |
| 59 | 157 Patrick | Sully | (null) |
| 60 | 158 Allan | McEwen | (null) |
| 61 | 159 Lindsey | Smith | (null) |
| 62 | 160 Louise | Doran | (null) |
| 63 | 161 Sarath | Sewall | (null) |
| 64 | 162 Clara | Vishney | (null) |
| 65 | 163 Danielle | Greene | (null) |
| 66 | 164 Mattea | Marvins | (null) |
| 67 | 165 David | Lee | (null) |
| 68 | 166 Sundar | Ande | (null) |
| 69 | 167 Amit | Banda | (null) |
| 70 | 168 Lisa | Ozer | (null) |
| 71 | 169 Harrison | Bloom | (null) |
| 72 | 170 Tayler | Fox | (null) |
| 73 | 171 William | Smith | (null) |
| 74 | 172 Elizabeth | Bates | (null) |
| 75 | 173 Sundita | Kumar | (null) |
| 76 | 174 Ellen | Abel | (null) |
| 77 | 175 Alyssa | Hutton | (null) |
| 78 | 176 Jonathon | Taylor | SA_MAN |
| 79 | 176 Jonathon | Taylor | SA_REP |
| 80 | 177 Jack | Livingston | (null) |
| 81 | 178 Kimberly | Grant | (null) |
| 82 | 179 Charles | Johnson | (null) |
| 83 | 180 Winston | Taylor | (null) |
| 84 | 181 Jean | Fleaur | (null) |
| 85 | 182 Martha | Sullivan | (null) |
| 86 | 183 Girard | Geoni | (null) |
| 87 | 184 Nandita | Sarchand | (null) |
| 88 | 185 Alexis | Bull | (null) |
| 89 | 186 Julia | Dellinger | (null) |
| 90 | 187 Anthony | Cabrio | (null) |
| 91 | 188 Kelly | Chung | (null) |
| 92 | 189 Jennifer | Dilly | (null) |
| 93 | 190 Timothy | Gates | (null) |
| 94 | 191 Randall | Perkins | (null) |
| 95 | 192 Sarah | Bell | (null) |
| 96 | 193 Britney | Everett | (null) |
| 97 | 194 Samuel | McCain | (null) |
| 98 | 195 Vance | Jones | (null) |

| | | | |
|-----|--------------|-----------|------------|
| 99 | 196 Alana | Walsh | (null) |
| 100 | 197 Kevin | Feeney | (null) |
| 101 | 198 Donald | OConnell | (null) |
| 102 | 199 Douglas | Grant | (null) |
| 103 | 200 Jennifer | Whalen | AC_ACCOUNT |
| 104 | 200 Jennifer | Whalen | AD_ASST |
| 105 | 201 Michael | Hartstein | MK_REP |
| 106 | 202 Pat | Fay | (null) |
| 107 | 203 Susan | Mavris | (null) |
| 108 | 204 Hermann | Baer | (null) |
| 109 | 205 Shelley | Higgins | (null) |
| 110 | 206 William | Gietz | (null) |

110 rows selected.

Exercise Manual

- Joining the `jobs` and `job_history` tables, display the job title and employee ID for all jobs, whether or not they have job history.

| JOB_TITLE | EMPLOYEE_ID |
|------------------------------------|-------------|
| 1 Public Accountant | 101 |
| 2 Public Accountant | 200 |
| 3 Accounting Manager | 101 |
| 4 Administration Assistant | 200 |
| 5 President | (null) |
| 6 Administration Vice President | (null) |
| 7 Accountant | (null) |
| 8 Finance Manager | (null) |
| 9 Human Resources Representative | (null) |
| 10 Programmer | 102 |
| 11 Marketing Manager | (null) |
| 12 Marketing Representative | 201 |
| 13 Public Relations Representative | (null) |
| 14 Purchasing Clerk | (null) |
| 15 Purchasing Manager | (null) |
| 16 Sales Manager | 176 |
| 17 Sales Representative | 176 |
| 18 Shipping Clerk | (null) |
| 19 Stock Clerk | 114 |
| 20 Stock Clerk | 122 |
| 21 Stock Manager | (null) |

- Modify the above query to restrict the result set to jobs whose minimum salary exceeds 9000.

| JOB_TITLE | EMPLOYEE_ID |
|---------------------------------|-------------|
| 1 President | (null) |
| 2 Administration Vice President | (null) |
| 3 Sales Manager | 176 |

Exercise Manual

- Joining the `jobs` and `job_history` tables, display the job title, employee ID, and starting date for all employees who started in that job after Jan. 1, 1998. Include jobs **even if** they do not have any history.

| JOB_TITLE | EMPLOYEE_ID | START_DATE |
|------------------------------------|-------------|------------|
| 1 Public Accountant | 200 | 01-JUL-02 |
| 2 Accounting Manager | 101 | 28-OCT-01 |
| 3 Administration Assistant | (null) | (null) |
| 4 President | (null) | (null) |
| 5 Administration Vice President | (null) | (null) |
| 6 Accountant | (null) | (null) |
| 7 Finance Manager | (null) | (null) |
| 8 Human Resources Representative | (null) | (null) |
| 9 Programmer | 102 | 13-JAN-01 |
| 10 Marketing Manager | (null) | (null) |
| 11 Marketing Representative | 201 | 17-FEB-04 |
| 12 Public Relations Representative | (null) | (null) |
| 13 Purchasing Clerk | (null) | (null) |
| 14 Purchasing Manager | (null) | (null) |
| 15 Sales Manager | 176 | 01-JAN-07 |
| 16 Sales Representative | 176 | 24-MAR-06 |
| 17 Shipping Clerk | (null) | (null) |
| 18 Stock Clerk | 114 | 24-MAR-06 |
| 19 Stock Clerk | 122 | 01-JAN-07 |
| 20 Stock Manager | (null) | (null) |

- Modify the above query: remove the start date restriction and also include the employee's first and last names.

| JOB_TITLE | EMPLOYEE_ID | START_DATE | FIRST_NAME | LAST_NAME |
|------------------------------------|-------------|------------|------------|-----------|
| 1 Accounting Manager | 101 | 28-OCT-01 | Neena | Kochhar |
| 2 Public Accountant | 101 | 21-SEP-97 | Neena | Kochhar |
| 3 Programmer | 102 | 13-JAN-01 | Lex | De Haan |
| 4 Stock Clerk | 114 | 24-MAR-06 | Den | Raphaely |
| 5 Stock Clerk | 122 | 01-JAN-07 | Payam | Kaufling |
| 6 Sales Representative | 176 | 24-MAR-06 | Jonathon | Taylor |
| 7 Sales Manager | 176 | 01-JAN-07 | Jonathon | Taylor |
| 8 Administration Assistant | 200 | 17-SEP-95 | Jennifer | Whalen |
| 9 Public Accountant | 200 | 01-JUL-02 | Jennifer | Whalen |
| 10 Marketing Representative | 201 | 17-FEB-04 | Michael | Hartstein |
| 11 Marketing Manager | (null) | (null) | (null) | (null) |
| 12 Public Relations Representative | (null) | (null) | (null) | (null) |
| 13 Purchasing Clerk | (null) | (null) | (null) | (null) |
| 14 Human Resources Representative | (null) | (null) | (null) | (null) |
| 15 Accountant | (null) | (null) | (null) | (null) |
| 16 Administration Vice President | (null) | (null) | (null) | (null) |
| 17 Shipping Clerk | (null) | (null) | (null) | (null) |
| 18 President | (null) | (null) | (null) | (null) |
| 19 Stock Manager | (null) | (null) | (null) | (null) |
| 20 Finance Manager | (null) | (null) | (null) | (null) |
| 21 Purchasing Manager | (null) | (null) | (null) | (null) |

Bonus Exercise (if time permits)

- Joining the `employees`, `job_history`, and `jobs` tables, display the job title, employee ID, start date, and employee first and last names for ALL employees, whether or not they have any job history.

Hint: You will need to change the most important table.

| JOB_TITLE | EMPLOYEE_ID | START_DATE | FIRST_NAME | LAST_NAME |
|----------------------------|-------------|------------|------------|-----------|
| 1 Public Accountant | 101 | 21-SEP-97 | Neena | Kochhar |
| 2 Public Accountant | 200 | 01-JUL-02 | Jennifer | Whalen |
| 3 Accounting Manager | 101 | 28-OCT-01 | Neena | Kochhar |
| 4 Administration Assistant | 200 | 17-SEP-95 | Jennifer | Whalen |
| 5 Programmer | 102 | 13-JAN-01 | Lex | De Haan |
| 6 Marketing Representative | 201 | 17-FEB-04 | Michael | Hartstein |
| 7 Sales Manager | 176 | 01-JAN-07 | Jonathon | Taylor |
| 8 Sales Representative | 176 | 24-MAR-06 | Jonathon | Taylor |
| 9 Stock Clerk | 114 | 24-MAR-06 | Den | Raphaely |
| 10 Stock Clerk | 122 | 01-JAN-07 | Payam | Kaufling |
| 11 (null) | (null) | (null) | Sundita | Kumar |
| 12 (null) | (null) | (null) | Diana | Lorentz |
| 13 (null) | (null) | (null) | Nancy | Greenberg |
| 14 (null) | (null) | (null) | Kevin | Mourgos |
| 15 (null) | (null) | (null) | Sarath | Sewall |
| 16 (null) | (null) | (null) | Sundar | Ande |

| | | | | |
|------------|--------|--------|---------|-----------|
| 96 (null) | (null) | (null) | John | Seo |
| 97 (null) | (null) | (null) | Sarah | Bell |
| 98 (null) | (null) | (null) | Nanette | Cambrault |
| 99 (null) | (null) | (null) | Shelley | Higgins |
| 100 (null) | (null) | (null) | David | Lee |
| 101 (null) | (null) | (null) | Alyssa | Hutton |
| 102 (null) | (null) | (null) | Steven | King |
| 103 (null) | (null) | (null) | Shanta | Vollman |
| 104 (null) | (null) | (null) | Alberto | Errazuriz |
| 105 (null) | (null) | (null) | Valli | Pataballa |
| 106 (null) | (null) | (null) | Stephen | Stiles |
| 107 (null) | (null) | (null) | Peter | Tucker |
| 108 (null) | (null) | (null) | Louise | Doran |
| 109 (null) | (null) | (null) | Martha | Sullivan |
| 110 (null) | (null) | (null) | Sigal | Tobias |

110 rows selected.

Congratulations! You have completed this exercise.

Chapter 5: Additional SQL Functions

Exercise 5.1: Additional SQL Functions

Connect to the HR account.

- Write a query to display the department ID, first and last, names, hire date, and commission percentage of all employees with manager 100 (Steven King) whose hire date is within 2 years of Jan 1, 2007. Sequence the list in department, hire date order. Use a non-standard function to display null commissions as 0.

| DEPARTMENT_ID | FIRST_NAME | LAST_NAME | HIRE_DATE | COMMISSION |
|---------------|------------|-----------|-------------|------------|
| 50 | Adam | Fripp | 10-APR-2005 | 0 |
| 50 | Shanta | Vollman | 10-OCT-2005 | 0 |
| 50 | Kevin | Mourgos | 16-NOV-2007 | 0 |
| 80 | Karen | Partners | 05-JAN-2005 | .3 |
| 80 | Alberto | Errazuriz | 10-MAR-2005 | .3 |
| 80 | Gerald | Cambrault | 15-OCT-2007 | .3 |
| 80 | Eleni | Zlotkey | 29-JAN-2008 | .2 |
| 90 | Neena | Kochhar | 21-SEP-2005 | 0 |

- Change the previous query to use a standard function to display the commission percentage. This time, order the list so that those hired closest to Jan 1, 2007, are listed first.

| DEPARTMENT_ID | FIRST_NAME | LAST_NAME | HIRE_DATE | COMMISSION |
|---------------|------------|-----------|-------------|------------|
| 80 | Gerald | Cambrault | 15-OCT-2007 | .3 |
| 50 | Kevin | Mourgos | 16-NOV-2007 | 0 |
| 80 | Eleni | Zlotkey | 29-JAN-2008 | .2 |
| 50 | Shanta | Vollman | 10-OCT-2005 | 0 |
| 90 | Neena | Kochhar | 21-SEP-2005 | 0 |
| 50 | Adam | Fripp | 10-APR-2005 | 0 |
| 80 | Alberto | Errazuriz | 10-MAR-2005 | .3 |
| 80 | Karen | Partners | 05-JAN-2005 | .3 |

Congratulations! You have completed this exercise.

This page intentionally left blank.

Chapter 6: Data Manipulation Language

Exercise 6.1: Manipulating Data

Connect to the HR account.

1. Display all the rows in the `regions` table.
2. Add a new row for Central America. Make it ID 5.
3. Display all the rows in the `regions` table.
4. Add a new row for South America. Make it ID 6.
5. Display all the rows in the `regions` table.
6. Update all `regions` rows with the name, Central America. Change their name to South and Central America.
7. Display all the rows in the `regions` table.
8. Delete the `regions` row whose ID is 6.
9. Display all the rows in the `regions` table.
10. Issue a ROLLBACK and re-display the `regions` table.

Congratulations! You have completed this exercise.

This page intentionally left blank.

Chapter 7: Databases with JDBC (Java Database Connectivity)

Exercise 7.1: Connecting to a Database

In this exercise, you will write the Java code to connect to the database.

Over the next two days, you will enhance this project to use and test a Data Access Object (DAO), a `TransactionManager`, and a Business Service.

This exercise will complete the `SimpleDataSource` class that provides a database Connection for the DAO to use. In a real-world project, you will not be creating a `DataSource` yourself. In most situations, an application server will create a `DataSource` and make it available for application programmers to use.

For the most part, the details of our `SimpleDataSource` are not that important. The `SimpleDataSource` is not part of the actual application source code. Instead, it is in the testing section of the Eclipse project because it will only be used during testing.

To make the `SimpleDataSource` functional, you will implement the `openConnection()` method. The instructions below will show you how to do this in a Test-Driven Development (TDD) style.

1. Open the `EmployeeManagement` project in Eclipse.
 - a. Verify that the Oracle dependency is listed in `pom.xml`.

```
<dependency>
  <groupId>com.oracle.database.jdbc</groupId>
  <artifactId>ojdbc8</artifactId>
  <version>18.3.0.0</version>
</dependency>
```
2. Open the `SimpleDataSourceTest.java` file.
 - a. It is located in the `com.fidelity.integration` package in `src/test/java`.
3. Verify the `@AfterEach` method calls the `shutdown` method on the `SimpleDataSource` object.
 - a. This ensures that the database Connection is closed after each test.
4. Run the test in this file.
 - a. You should get a **Green** bar.
 - b. This shows that the `SimpleDataSource` object is successfully created.

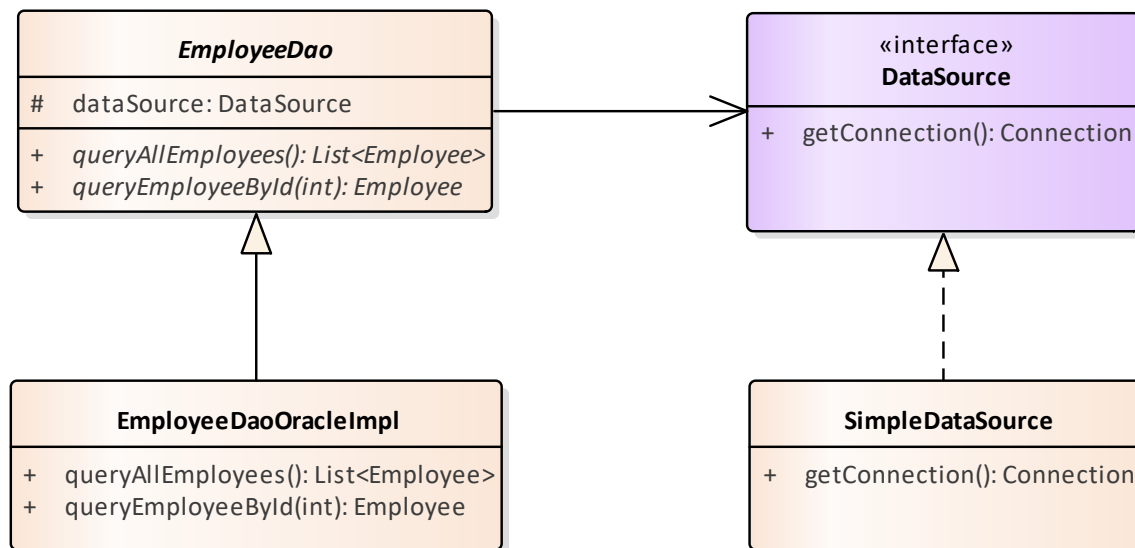
5. Open the `SimpleDataSource.java` file.
 - a. In the `com.fidelity.integration` package in `src/test/java`.
6. Examine the `openConnection()` method.
 - a. You will notice that the method is not entirely functional at this point.
 - b. The `openConnection` method is private.
 - c. The `getConnection` method calls the `openConnection` method.
7. In `SimpleDataSourceTest.java`, write a JUnit test to get a `Connection` from the `SimpleDataSource`.
 - a. By calling the `getConnection` method.
 - b. Assert that the `Connection` that is returned is not null.
 - c. Run the test.
 - d. Naturally, the test will fail!
 - e. But of course, you expected that.
 - f. **Red** bar, ...
8. Complete the `openConnection` method to return a `Connection` to the Oracle database.
 - a. Review the file `db.properties`.
 - i. This file is in the `src/main/resources` folder.
 - ii. It will be in the classpath at runtime.
 - b. Call `DriverManager.getConnection` to create the connection.
 - iii. Start with the simplest code that could possibly work. See the example in the course notes that defines the connection parameters as hard-coded strings.
 - iv. After you get the **Green** bar for your test case, refactor your `getConnection` method to read the connection parameters from a properties file. See the second example in the course notes.
 - c. If an exception occurs, throw a `DatabaseException`.
9. Run your test again.
 - a. Of course, you now get a **Green** bar because you wrote perfect code!
 - b. **Red** bar, **Green** bar, ...

Congratulations! You have completed this exercise.

Exercise 7.2: Creating Objects from a Database Query

In this exercise, you will perform queries to access employee data from the `scott.emp` table.

The UML class diagram for this exercise is shown below:



- Open the SQL Developer application.
 - View the structure of the `emp` table in the `scott` database.
 - In a worksheet, write a query that will select all of the records in the `emp` table.
 - Do **not** use `SELECT *`.
 - List the columns by name in the `SELECT` statement.
- Continue working with the `EmployeeManagement` project in Eclipse.
- Examine the `Employee` class.
 - This is the Java representation of the `scott.emp` table.
- Examine the `EmployeeDao.java` file.
 - Note that the class is abstract.
 - The two query methods are also abstract.
 - There is a protected `DataSource` field.

5. Create a new Java class in the `com.fidelity.integration` package to hold the database access code.
 - a. Name it `EmployeeDaoOracleImpl`
 - b. Have it extend the `EmployeeDao` base class.
 - c. Put it in the `com.fidelity.integration` package.
 - i. In all the solution projects, we will use `com.fidelity.integration` to show that it is our integration layer.
 - d. Let Eclipse write the two abstract methods inherited from the base class.
6. Now, create a test class for the `EmployeeDaoOracleImpl`.
 - a. Right-click `EmployeeDaoOracleImpl` and select **New | Other...**, then navigate to **Java | Junit** and select **Junit Test Case**.
 - b. Make sure you are creating a JUnit 5 (Jupiter) test case.
 - c. Check that the name of the class is `EmployeeDaoQueryTest`, and it is in the same package as `EmployeeDaoOracleImpl`.
 - d. Change the Source folder to `EmployeeManagement/src/test/java`.
7. Check the `@BeforeEach` and `@AfterEach` checkboxes.
8. In the `@BeforeEach` method, create a `SimpleDataSource` and an `EmployeeDaoOracleImpl` object.
9. In the `@AfterEach` methods, call the `shutdown()` method on the `SimpleDataSource` object.
10. Write a Test method to retrieve all records from the `scott.emp` table and return a list of `Employee` objects.
 - a. Work in TDD fashion.
 - b. How will you tell whether the method returned the right set of employees?
11. Run the test.
 - a. You got the **Red** bar, didn't you?
12. Write the implementation of the method in `EmployeeDaoOracleImpl`.
 - a. Use the sample code in the course notes as a guide for writing this new method.
 - b. Use the query that you created in SQL Developer.
 - i. Remember to delete the trailing ";"
 - c. Of course, you will use a `PreparedStatement`!

- d. Although they are not the only nullable columns, both `mgr` and `comm` have null values. JDBC will treat these as 0, which is OK for now.
13. Run the test
- a. Got the **Green** bar? Great! Go to Step 14.
 - b. Got a **Red** bar?
 - i. Read the error message.
 - ii. Do some debugging.
 - iii. Get that coveted **Green** bar.
14. Create a method to retrieve a single `scott.emp` record by `empno` and return a single `Employee` object.
- a. Again, work in TDD fashion. What tests do you need?
 - b. Write the query in SQL Developer first, then copy into the DAO method.
 - i. If a SQL statement does not work in SQL Developer, it will definitely not work in your JDBC code!
 - c. Write the tests, and then let Eclipse create the method for you.
15. Work until you get that **Green** bar.

Optional steps, if you have time:

16. Create a method to retrieve all the `scott.emp` records where the employee name matches a string that is passed in as a parameter. Return a list of `Employee` objects.
- a. Even though all the names are actually unique, this is not guaranteed by the database, so you should return a list, not a single `Employee` object.
17. Add a test that proves that SQL injection is not possible. Consider the string `"BobbyTables" OR '1' = '1'`.
18. Create a method to retrieve all the `scott.emp` records in a particular department and return a list of `Employee` objects.

Congratulations! You have completed this exercise.

This page was intentionally left blank.

Chapter 8: Updating Databases

Exercise 8.1: Inserting a Record with JDBC

Since the `insertEmployee` method will insert a new record into the database, for testing, we should run the test in a transaction and roll back the transaction after the test has completed.

If you make significant changes to the `scott.emp` table, you can use the file `sql\emp_populate.sql` to recover the table.

In this exercise, you will use a `TransactionManager` to start and roll back a transaction.

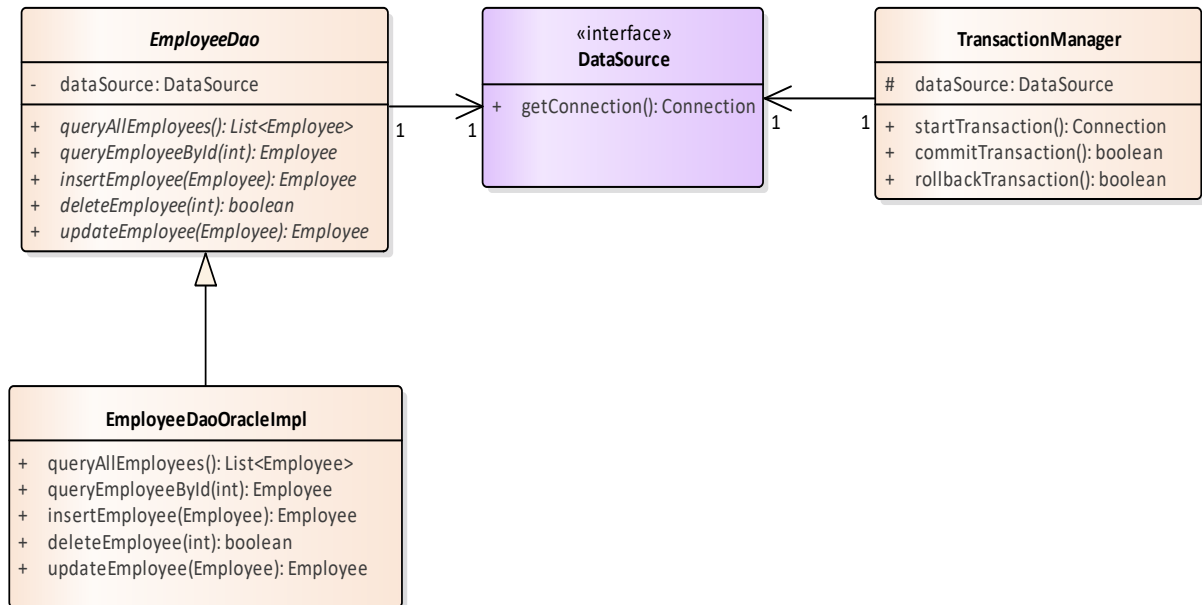
The `TransactionManager` class lives up to its name by starting, committing, and rolling back transactions.

The DAO class will not call on (or even know of) the `TransactionManager`.

1. Create a new class `TransactionManager` in the `com.fidelity.integration` package.
2. Define a constructor for `TransactionManager` that has a `DataSource` argument and initializes a data field.
3. Define the `startTransaction`, `commitTransaction`, and `rollbackTransaction` methods.
 - a. Use the `Connection` from the `DataSource` in the transaction methods.
4. Open the `EmployeeDaoDMLTest.java` file.
5. In the `@BeforeEach` method, create a new `EmployeeDaoOracleImpl` object and assign it to the DAO data member.
6. Call the `TransactionManager startTransaction()` in `@BeforeEach`.
7. Call the `TransactionManager rollbackTransaction()` in `@AfterEach`.
8. Add the `insertEmployee` method to the `EmployeeDao` base class.
9. Write a test for the `insertEmployee` method in `EmployeeDaoDMLTest`.
10. Let Eclipse write the method in `EmployeeDaoOracleImpl`.
11. Run the test to get the **Red** bar.
12. Define the new method correctly in `EmployeeDaoOracleImpl`.

Exercise Manual

13. Run the test to get the **Green** bar.

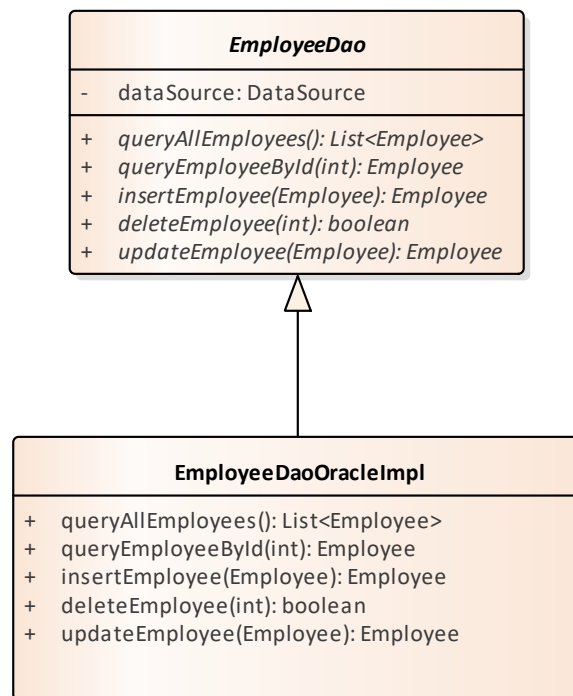


Congratulations! You have completed this exercise.

Exercise 8.2: Updating and Deleting Records

In this exercise, you will continue your work on the `EmployeeManagement_DML_Test` class and the `EmployeeDaoOracleImpl` class.

1. Add the abstract `updateEmployee` method to the `EmployeeDao` class.
2. Write a test for the new method.
3. Implement the new method in the `EmployeeDaoOracleImpl`.
4. Get that **Green** bar.
5. Add the `deleteEmployee` method to the `EmployeeDao` class.
6. Write a test for the new method.
7. Implement the new method in the `EmployeeDaoOracleImpl`.
8. Get that **Green** bar.
9. Write a negative test to verify the DAO throws a `DatabaseException`.
 - a. Attempt to update an employee to be in a non-existent department.
 - b. Verify the `DatabaseException` is thrown.



Congratulations! You have completed this exercise.

Exercise 8.3: Working with Multi-Table Queries

In this exercise, you will extract information from two database tables.

In this scenario, a query will extract data from two tables. One record in each table will be involved in the query. A record in the `emp` table will be associated with one record in the `emp_perf` table.

1. Use SQL Developer to open the `emp_perf.sql` file that is in the `sql` folder in the `EmployeeManagement` Eclipse project.
2. Run the `emp_perf.sql` script.
 - a. This will create a new table named `emp_perf`.
 - b. Each record in the `emp_perf` table stores performance evaluation information for an employee.
3. Use a new worksheet in SQL Developer to create a query that will select all the columns in a record in the `emp` table and the `perf_code` column in the `emp_perf` table for that employee.
 - a. Use a `JOIN` to obtain the data from the two tables.
 - b. Use a `WHERE` clause to specify the employee number (`empno` column).
4. In Eclipse, create an `enum` with values for an employee performance review. When you process the query's result set, you will need to map the integer review code to a value of your new `enum`.
 - a. *Hint:* see the `emp_perf` table for the code values.
 - b. *Hint:* refer to the Enumerated Types section in Chapter 7 of the course notes.
5. Modify the `Employee` class to contain a property to store the performance review value for that employee.
 - a. This property can be null if there is no record for that employee in the `emp_perf` table.
 - b. Create the getter and setter for the new property.
 - c. Update the `equals()`, `hashCode()`, and `toString()` methods.
 - i. Ask Eclipse to do this.
6. Using TDD, define and test a new DAO method that queries for an `Employee` by `empno` including the performance property.
 - a. Set the performance property for the employee to the value in the `emp_perf` table for that employee.
 - b. Handle the situation when the employee does not have an entry in the `emp_perf` table.

Congratulations! You have completed this exercise.

Chapter 9: Working with a Data Access Object

Exercise 9.1: Debugging Data Access Object Methods

Time: 20 minutes

Format: Individual programming exercise

Once again, it's Friday afternoon and you are ready to pack up for a well-deserved weekend after your hard work developing a Data Access Object that utilizes JDBC.

Life is good! The weekend is almost here.

Before your weekend can start, there are a few JUnit tests that have to be examined. There are several tests, but there are two failures and one error.

Since you are now a seasoned Java and JDBC developer and debugger, you are not panicking. You know what you have to do to track the problems down and eliminate them ASAP.

1. Examine the code in the `ProblematicEmployee` project that defines the Data Access Object that is communicating with the `scott` database.
2. Run the JUnit tests for the project.
3. Notice that you get some **Green** bars but also a nasty **Red** bar indicating one error and two failures.
4. Dust off your debugging skills and get to work!
 - a. Correct one problem at a time.
 - b. You may want to use the Debugger to step through the code in order to isolate the problem.
5. Now that you have all **Green**, you are feeling pretty good.
6. Are you completely sure that your DAO code is correct?
7. Examine the set of DAO tests.
8. Are there any other tests that should be written for this DAO?
 - a. If so, you know what to do.
 - b. **Red bar**, **Green bar**, **Refactor**.

Congratulations! You have completed this exercise.

Optional Exercise 9.2: Testing a Business Service with a Mock DAO

In this exercise, you will continue working with the `EmployeeManagement` project. You will define the `EmployeeManagementService` class and test with a mock DAO.

Now that you have a working (and tested) DAO, it is time to concentrate on how to use the DAO functionality in an application. Introducing a Business Service is a step in that direction.

The methods in the Business Service will correspond to requirements for the application that is being developed. The Business Service methods will call on the DAO to communicate with the back-end database.

During unit testing of the Business Service, we do not want to actually modify the contents of the database. For this reason, we will use Mockito to create a mock DAO that has the same interface as the `EmployeeDao` but does not actually communicate with the database. Mockito makes it easy to test edge conditions and error paths. You'll add service test cases that verify its behavior for an empty employee list and a DAO exception, as well as a positive test case for a non-empty list.

As usual, work using TDD: write a test case first, run it and watch it fail, then modify the method under test to make the test case pass.

1. The `EmployeeManagementService` is a business service that has a dependency on the `EmployeeDao`.
2. Modify `EmployeeManagementServiceTest` so it creates a mock DAO before each test case.
 - a. Refer to the examples in the course notes.
3. Add a new test case that verifies that, if the DAO returns a non-empty list of employees, the `EmployeeManagementService` `getAllEmployees()` method returns the same list.
 - a. After writing the test case, add a `getAllEmployees()` method to `EmployeeManagementService`.
 - b. Confirm the new unit test now passes.
4. Add a unit test that verifies that, if the DAO `getAllEmployees()` method returns an empty list of employees, the `EmployeeManagementService` `getAllEmployees()` method throws an `IllegalStateException`. Run the test case, watch it fail, then add the `EmployeeManagementService` `getAllEmployees()` method to make the test case pass.

5. Add a unit test that verifies that, if the DAO `getAllEmployees()` method throws a `DatabaseException`, the `getAllEmployees()` method throws the same exception. Run the test case, watch it fail, then modify `getAllEmployees()` to make the test case pass.

Bonus Exercise (if time permits)

6. A best practice is not to "leak" details about a class's implementation to its clients. Currently, if the DAO throws a `DatabaseException`, `EmployeeManagementService` passes the same exception back to its client, thus leaking details about its implementation (i.e., the fact that the service uses a database instead of a web service call). Now you'll fix this deficiency.
 - a. Define a new `ServiceException` class, which extends `RuntimeException`.
 - b. Add a unit test that verifies that, if the DAO `getAllEmployees()` method throws an exception, the `EmployeeManagementService` `getAllEmployees()` method throws the new `ServiceException`. Run the test case and watch it fail.
 - c. Modify `EmployeeManagementService` `getAllEmployees()` to make the test case pass.

Congratulations! You have completed this exercise.

Exercise 9.3 Putting It All Together

In this exercise, you will create a new project that will contain a Data Access Object that will define methods for queries, inserts, updates, and deletes on an Oracle database.

You will define tests for all of the DAO methods. Of course, you will do this TDD style!

The tests for the DML operations will be performed in a transaction that is rolled back after the test completes.

Your instructor will provide instructions on how to import the starter project into Eclipse.

Congratulations! You have completed this exercise.

Chapter 10: Advanced JDBC

Optional Exercise 10.1: JDBC Code Review

You have been asked to participate in a code review of a Data Access Object that will be used in airline flight search software. The software will run in every airport check-in kiosk for a major airline.

Open the `FlightTracker` project in Eclipse.

Work in small groups to examine the code in the `FlightSearch.java` source code file.

Feel free to run the JUnit tests that are defined in the `FlightSearchTest.java` file to verify the small set of unit tests pass and produce the coveted **Green** bar.

A spokesperson for your group should be prepared to discuss the group's evaluation of the `FlightSearch` code.

In particular, your evaluation should include any positive aspects of the code (assuming there are some), and any negative aspects of the code (should there be any).

Your group should decide if you would give the `FlightSearch` code the go-ahead for inclusion in the production software installed in the airport kiosks.

If you are not giving approval of the code, then be specific as to why you would not approve of the code and what must be corrected to gain approval.

Congratulations! You have completed this exercise.

Optional Exercise 10.2: One-to-Many Queries

In this exercise, you will once again extract data from two tables. There are two scenarios described below. In both scenarios, one `Department` object will have a collection of `Employee` objects.

One of the main challenges in this exercise is to determine how to properly process the `ResultSet` that is returned from the query.

Scenario 1: One-to-many (1-1..*)

In this scenario, a query will extract data from two tables. One record in the `dept` table will be associated with one to several records in the `emp` table.

1. Use SQL Developer to inspect the structure of the `dept` table in the `scott` database.
2. In your Eclipse project, create a `Department` class to store the data from a record in the `dept` table.
3. In the `Department` class, add a data field to store the `Employees` in that `Department`.
4. Using TDD, write and test a DAO method to query for a `Department` by id.
 - a. Return a `Department` object with the collection of `Employees` for that `Department`.
 - b. Use SQL Developer to help you write the query.

Scenario 2: One-to-many (1-0..*)

In this scenario, a query will extract data from two tables. One record in the `dept` table will be associated with zero to several records in the `emp` table. All of the records in the `dept` table are to be returned in this scenario, even if a `dept` record is not associated with any `emp` records.

5. Using TDD, write and test a DAO method that will query for all `Departments`.
 - a. Return a collection of `Department` objects with `Employees`.
 - b. Include `Departments` that do not contain any `Employees`.
 - c. Use SQL Developer to help you write the query.

Congratulations! You have completed this exercise.

Optional Exercise 10.3: Testing a Service Proxy with Mockito

The problem is to guarantee that the DML operations provided by the DAO are called in a transaction.

While the `EmployeeManagementService` could call the `startTransaction` and `commitTransaction` methods to ensure the DML operations are performed in a transaction, that seems to leak some implementation details from the DAO into the service.

An alternative is to use a proxy for the service. This gives us the opportunity to discuss design patterns again and to demonstrate the use of one in a “real” situation.

We want the proxy to call `startTransaction` on the `TransactionManager`, then call the business service methods that are to be part of the transaction, then call `commitTransaction` on the `TransactionManager`.

The `insertNewManager()` method in the service class inserts a new `Manager` and updates the `Employees` to be managed by the new `Manager`.

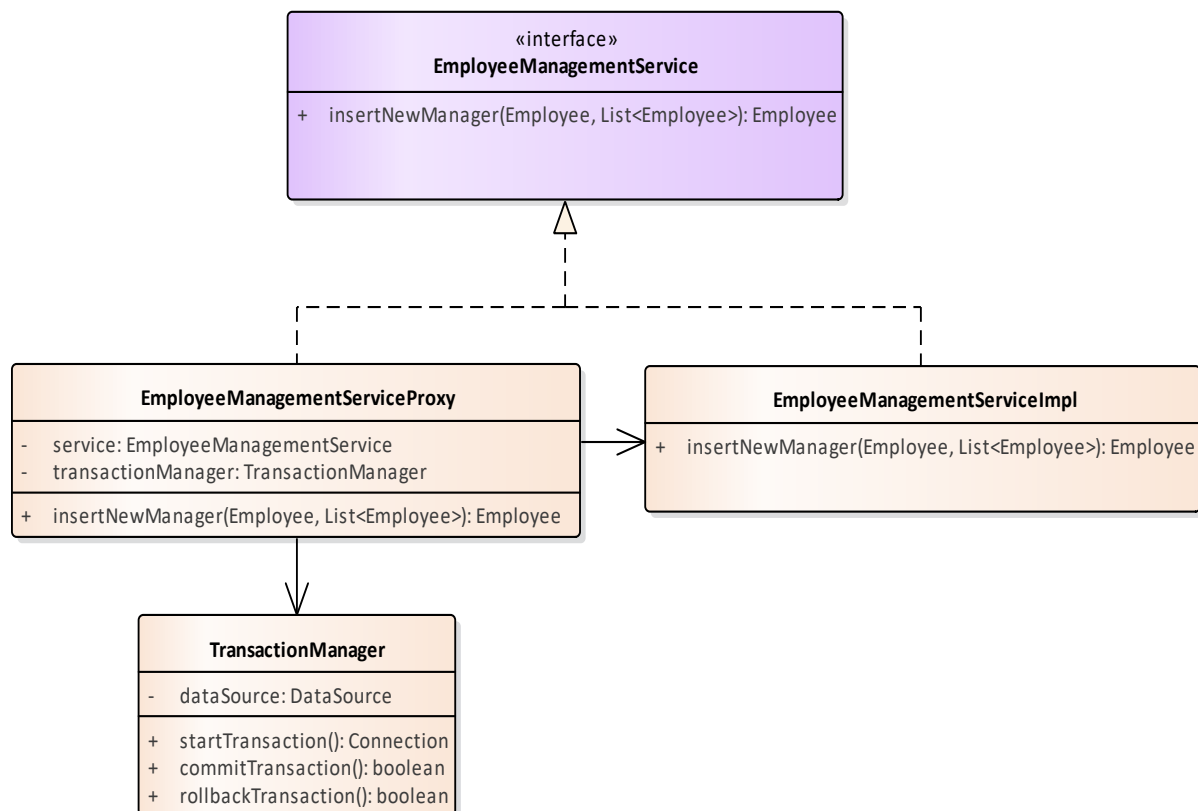
This method definitely requires a transaction to guarantee that all operations complete successfully, or all are rolled back.

If a `DatabaseException` is thrown, the proxy can let it pass on through to the client.

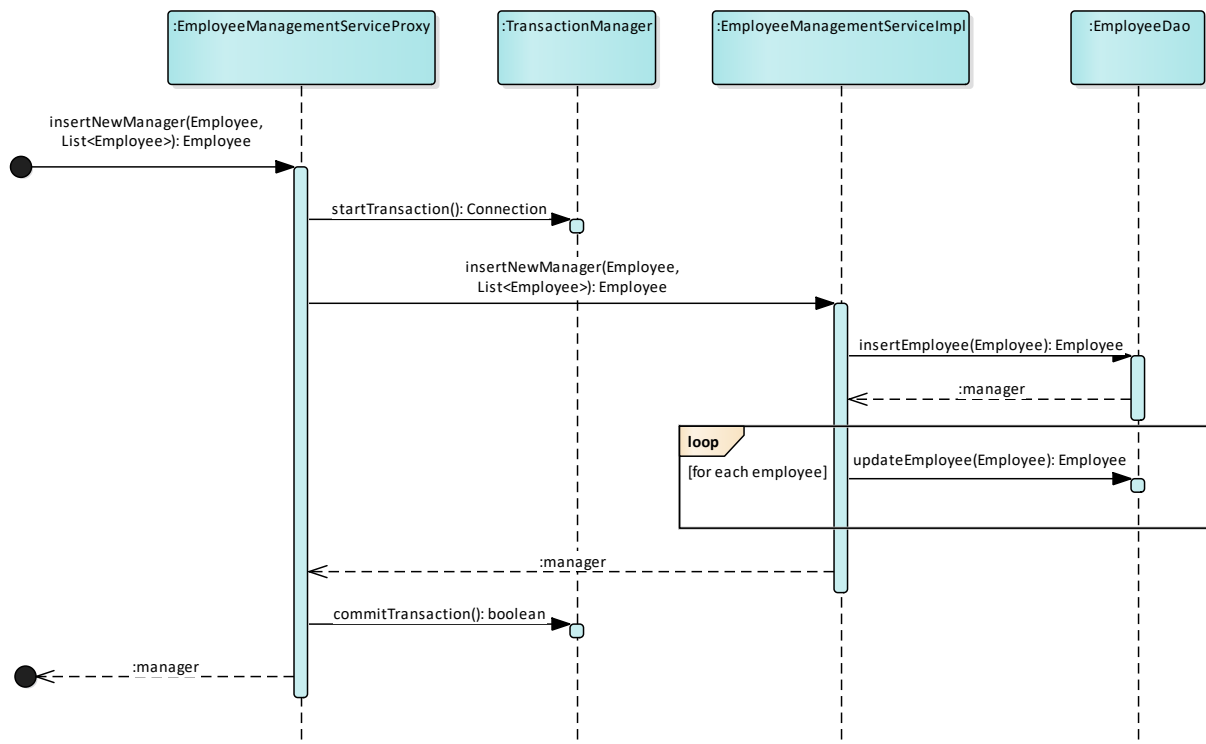
1. Rename the `EmployeeManagementService` class to be `EmployeeManagementServiceImpl`.
 - a. Use the **Refactor | Rename** support in Eclipse.
2. Run the tests to verify that all the tests pass.
 - a. That is what should happen when you apply Refactoring.
3. Extract the `EmployeeManagementService` interface from `EmployeeManagementServiceImpl`.
 - a. Use the **Refactor | Rename** support in Eclipse.
4. Run the tests to see that **Green** bar!
5. Create a new class named `EmployeeManagementServiceProxy`.
 - d. The `EmployeeManagementServiceProxy` class implements the `EmployeeManagementService` interface.
 - e. The `EmployeeManagementServiceProxy` class has an association with `EmployeeManagementServiceImpl`.

Exercise Manual

6. Open the `EmployeeManagementServiceProxyTest.java` file.
7. Uncomment the lines of code that:
 - a. Add a data field for an `EmployeeManagementServiceProxy`.
 - b. Add a data field for a `TransactionManager`.
8. Uncomment the lines of code in the `@BeforeEach` method that create Mockito mocks for:
 - a. `TransactionManager`
 - b. `EmployeeManagementServiceProxy`
9. Write a test to verify the `startTransaction`, `commitTransaction` methods are called on the `TransactionManager` when `insertNewManager` is called on the proxy.
 - a. Use Mockito to do the verification.
 - b. Use Mockito to verify the `insertNewManager` method is called on the `EmployeeManagementServiceImpl` when `insertNewManager` is called on the proxy.



Exercise Manual



Congratulations! You have completed this exercise.

Optional Exercise 10.4: Writing an Integration Test for the Service and DAO

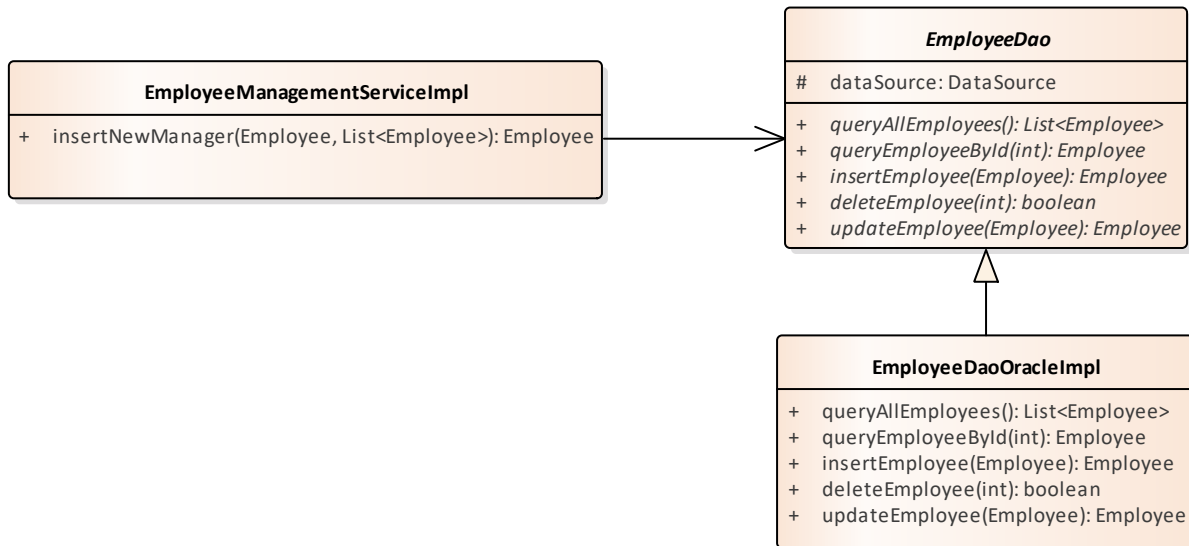
Since the `EmployeeManagementServiceImpl` class does not start a transaction directly, it is possible to test the insert, update, and delete operations and then roll the changes back in the `@AfterEach` method as you have done before.

Be sure to start a transaction in the `@BeforeEach` method and to roll back the transaction in the `@AfterEach` method.

You can start and roll back the transaction as was done in the DML tests.

1. Open the `ServiceDaoIntegrationTest.java` file.
 - a. Declare a `TransactionManager` field.
2. Do the following in the `@BeforeEach` method:
 - a. Create an `EmployeeDaoOracleImpl` and assign to the `dao` field.
 - b. Create an `EmployeeManagementService` and assign to the `service` field.
 - c. Add a `TransactionManager` field to the test class, then create a `TransactionManager` and assign to the `TransactionManager` field.
 - d. Use the `TransactionManager` to start a transaction.
3. Do the following in the `@AfterEach` method:
 - a. Use the `TransactionManager` to roll back the transaction.
4. Use TDD to implement the `insertNewManager()` method. First, write an integration test case.
 - f. `insertNewManager()` takes two arguments: the manager (an `Employee` instance) and a list of `Employees` who report to the manager.
 - g. The test case should verify the database contents before and after the call to `insertNewManager()`.
5. `EmployeeManagementService.insertNewManager()` should call methods of the `EmployeeDao` to perform the necessary database operations.
 - a. Insert the new manager.
 - b. Insert all the manager's employees.
 - c. If the DAO throws an exception during any database operation, `insertNewManager()` should throw a `ServiceException`.

6. Now run the integration test and verify that you get a green bar.



Congratulations! You have completed this exercise.

This page was intentionally left blank.

Chapter 11: Aggregating Information

Exercise 11.1: Using the Aggregate Functions

Connect to the SCOTT account.

- Write a query displaying how many rows there are in the emp table.

```
Count
-----
14
```

- Write a query displaying the empno, name, salary, and commission for all rows in the emp table, sequencing the list in salary (ascending) order.

| EMPNO | ENAME | SAL | COMM |
|-------|--------|------|------|
| 7369 | SMITH | 800 | |
| 7900 | JAMES | 950 | |
| 7876 | ADAMS | 1100 | |
| 7521 | WARD | 1250 | 500 |
| 7654 | MARTIN | 1250 | 1400 |
| 7934 | MILLER | 1300 | |
| 7844 | TURNER | 1500 | 0 |
| 7499 | ALLEN | 1600 | 300 |
| 7782 | CLARK | 2450 | |
| 7698 | BLAKE | 2850 | |
| 7566 | JONES | 2975 | |
| 7788 | SCOTT | 3000 | |
| 7902 | FORD | 3000 | |
| 7839 | KING | 5000 | |

- Write a query displaying how many non-null salary values exist and how many distinct non-null salary values exist in the emp table.

| Count | CDistinct |
|-------|-----------|
| 14 | 12 |

- Write a query displaying how many non-null commission values exists, the sum of the non-null commission values, the average of the non-null commission values for all rows in the emp table.

| Count | Sum | Average |
|-------|------|---------|
| 4 | 2200 | 550 |

- Modify the above query by adding the average of commission values treating unknown values as zero. Round this value to three decimal places.

| Count | Sum | Average | Average of all Records |
|-------|------|---------|------------------------|
| 4 | 2200 | 550 | 157.143 |

Exercise Manual

6. Write a query displaying the largest and smallest salaries in the `emp` table.

| Maximum Salary | Minimum Salary |
|----------------|----------------|
| ----- | ----- |
| 5000 | 800 |

7. Write a query displaying the latest and the earliest hire dates in the `emp` table.

| Maximum Hire Date | Minimum Hire Date |
|-------------------|-------------------|
| ----- | ----- |
| 12-JAN-1983 | 17-DEC-1980 |

Congratulations! You have completed this exercise.

Exercise 11.2: GROUP BY and HAVING

Connect to the HR account.

- For each department in the `employees` table, show the total count of employees, the highest salary, the smallest salary, the sum of the salaries, and the average of salaries (round to the nearest whole currency unit).

| DEPARTMENT_ID | COUNT(*) | MIN(SALARY) | MAX(SALARY) | Total Salary | Avg Salary |
|---------------|----------|-------------|-------------|--------------|------------|
| 10 | 1 | 4400 | 4400 | 4400 | 4400 |
| 20 | 2 | 6000 | 13000 | 19000 | 9500 |
| 30 | 6 | 2500 | 11000 | 24900 | 4150 |
| 40 | 1 | 6500 | 6500 | 6500 | 6500 |
| 50 | 45 | 2100 | 8200 | 156400 | 3476 |
| 60 | 5 | 4200 | 9000 | 28800 | 5760 |
| 70 | 1 | 10000 | 10000 | 10000 | 10000 |
| 80 | 34 | 6100 | 14000 | 304500 | 8956 |
| 90 | 3 | 17000 | 24000 | 58000 | 19333 |
| 100 | 6 | 6900 | 12008 | 51608 | 8601 |
| 110 | 2 | 8300 | 12008 | 20308 | 10154 |
| | 1 | 7000 | 7000 | 7000 | 7000 |

Note: Your output may not be in the same sequence.

- Modify the presentation sequence of the above query: the departments should be in ascending average salary order.

| DEPARTMENT_ID | COUNT(*) | MIN(SALARY) | MAX(SALARY) | Total Salary | Avg Salary |
|---------------|----------|-------------|-------------|--------------|------------|
| 50 | 45 | 2100 | 8200 | 156400 | 3476 |
| 30 | 6 | 2500 | 11000 | 24900 | 4150 |
| 10 | 1 | 4400 | 4400 | 4400 | 4400 |
| 60 | 5 | 4200 | 9000 | 28800 | 5760 |
| 40 | 1 | 6500 | 6500 | 6500 | 6500 |
| | 1 | 7000 | 7000 | 7000 | 7000 |
| 100 | 6 | 6900 | 12008 | 51608 | 8601 |
| 80 | 34 | 6100 | 14000 | 304500 | 8956 |
| 20 | 2 | 6000 | 13000 | 19000 | 9500 |
| 70 | 1 | 10000 | 10000 | 10000 | 10000 |
| 110 | 2 | 8300 | 12008 | 20308 | 10154 |
| 90 | 3 | 17000 | 24000 | 58000 | 19333 |

- Modify the previous query by adding a new column: calculate how much each department's smallest salary is below the average salary. Sequence the list by this expression.

| DEPARTMENT_ID | COUNT(*) | MIN(SALARY) | MAX(SALARY) | Total Salary | Avg Salary | Below Avg |
|---------------|----------|-------------|-------------|--------------|------------|-----------|
| 20 | 2 | 6000 | 13000 | 19000 | 9500 | 3500 |
| 80 | 34 | 6100 | 14000 | 304500 | 8956 | 2856 |
| 90 | 3 | 17000 | 24000 | 58000 | 19333 | 2333 |
| 110 | 2 | 8300 | 12008 | 20308 | 10154 | 1854 |
| 100 | 6 | 6900 | 12008 | 51608 | 8601 | 1701 |
| 30 | 6 | 2500 | 11000 | 24900 | 4150 | 1650 |
| 60 | 5 | 4200 | 9000 | 28800 | 5760 | 1560 |
| 50 | 45 | 2100 | 8200 | 156400 | 3476 | 1376 |
| 40 | 1 | 6500 | 6500 | 6500 | 6500 | 0 |
| 70 | 1 | 10000 | 10000 | 10000 | 10000 | 0 |
| 10 | 1 | 4400 | 4400 | 4400 | 4400 | 0 |
| | 1 | 7000 | 7000 | 7000 | 7000 | 0 |

Exercise Manual

4. Modify the above query by changing the analysis: we now want to know all the above information by the manager each employee works for.

| MANAGER_ID | COUNT(*) | MIN(SALARY) | MAX(SALARY) | Total Salary | Avg Salary | Below Avg |
|------------|----------|-------------|-------------|--------------|------------|-----------|
| 100 | 14 | 5800 | 17000 | 155400 | 11100 | 5300 |
| 101 | 5 | 4400 | 12008 | 44916 | 8983 | 4583 |
| 148 | 6 | 6100 | 11500 | 51900 | 8650 | 2550 |
| 149 | 6 | 6200 | 11000 | 50000 | 8333 | 2133 |
| 147 | 6 | 6200 | 10500 | 46600 | 7767 | 1567 |
| 146 | 6 | 7000 | 10000 | 51000 | 8500 | 1500 |
| 145 | 6 | 7000 | 10000 | 51000 | 8500 | 1500 |
| 121 | 8 | 2100 | 4200 | 25400 | 3175 | 1075 |
| 108 | 5 | 6900 | 9000 | 39600 | 7920 | 1020 |
| 103 | 4 | 4200 | 6000 | 19800 | 4950 | 750 |
| 122 | 8 | 2200 | 3800 | 23600 | 2950 | 750 |
| 123 | 8 | 2500 | 4000 | 25900 | 3238 | 738 |
| 120 | 8 | 2200 | 3200 | 22100 | 2763 | 563 |
| 124 | 8 | 2500 | 3500 | 23000 | 2875 | 375 |
| 114 | 5 | 2500 | 3100 | 13900 | 2780 | 280 |
| 102 | 1 | 9000 | 9000 | 9000 | 9000 | 0 |
| | 1 | 24000 | 24000 | 24000 | 24000 | 0 |
| 201 | 1 | 6000 | 6000 | 6000 | 6000 | 0 |
| 205 | 1 | 8300 | 8300 | 8300 | 8300 | 0 |

5. Another analysis request has been made: modify the previous query to “rate” managers within each department by how far their lowest employee salary is below average.

| DEPTID | MGRID | COUNT(*) | MIN(SALARY) | MAX(SALARY) | Total Salary | Avg Salary | Below Avg |
|--------|-------|----------|-------------|-------------|--------------|------------|-----------|
| 80 | 148 | 6 | 6100 | 11500 | 51900 | 8650 | 2550 |
| 80 | 149 | 5 | 6200 | 11000 | 43000 | 8600 | 2400 |
| 80 | 100 | 5 | 10500 | 14000 | 61000 | 12200 | 1700 |
| 80 | 147 | 6 | 6200 | 10500 | 46600 | 7767 | 1567 |
| 80 | 146 | 6 | 7000 | 10000 | 51000 | 8500 | 1500 |
| 80 | 145 | 6 | 7000 | 10000 | 51000 | 8500 | 1500 |
| 50 | 100 | 5 | 5800 | 8200 | 36400 | 7280 | 1480 |
| 50 | 121 | 8 | 2100 | 4200 | 25400 | 3175 | 1075 |
| 100 | 108 | 5 | 6900 | 9000 | 39600 | 7920 | 1020 |
| 60 | 103 | 4 | 4200 | 6000 | 19800 | 4950 | 750 |
| 50 | 122 | 8 | 2200 | 3800 | 23600 | 2950 | 750 |
| 50 | 123 | 8 | 2500 | 4000 | 25900 | 3238 | 738 |
| 50 | 120 | 8 | 2200 | 3200 | 22100 | 2763 | 563 |
| 50 | 124 | 8 | 2500 | 3500 | 23000 | 2875 | 375 |
| 30 | 114 | 5 | 2500 | 3100 | 13900 | 2780 | 280 |
| 10 | 101 | 1 | 4400 | 4400 | 4400 | 4400 | 0 |
| 90 | | 1 | 24000 | 24000 | 24000 | 24000 | 0 |
| 30 | 100 | 1 | 11000 | 11000 | 11000 | 11000 | 0 |
| 110 | 205 | 1 | 8300 | 8300 | 8300 | 8300 | 0 |
| 60 | 102 | 1 | 9000 | 9000 | 9000 | 9000 | 0 |
| 100 | 101 | 1 | 12008 | 12008 | 12008 | 12008 | 0 |
| 90 | 100 | 2 | 17000 | 17000 | 34000 | 17000 | 0 |
| 20 | 100 | 1 | 13000 | 13000 | 13000 | 13000 | 0 |
| 20 | 201 | 1 | 6000 | 6000 | 6000 | 6000 | 0 |
| 110 | 101 | 1 | 12008 | 12008 | 12008 | 12008 | 0 |
| 70 | 101 | 1 | 10000 | 10000 | 10000 | 10000 | 0 |
| 40 | 101 | 1 | 6500 | 6500 | 6500 | 6500 | 0 |
| | 149 | 1 | 7000 | 7000 | 7000 | 7000 | 0 |

Exercise Manual

6. Modify the above query to show only those managers within a department that have more than 5 employees reporting to them.

| DEPTID | MGRID | COUNT (*) | MIN (SALARY) | MAX (SALARY) | Total Salary | Avg Salary | Below Avg |
|--------|-------|-----------|--------------|--------------|--------------|------------|-----------|
| 80 | 148 | 6 | 6100 | 11500 | 51900 | 8650 | 2550 |
| 80 | 147 | 6 | 6200 | 10500 | 46600 | 7767 | 1567 |
| 80 | 146 | 6 | 7000 | 10000 | 51000 | 8500 | 1500 |
| 80 | 145 | 6 | 7000 | 10000 | 51000 | 8500 | 1500 |
| 50 | 121 | 8 | 2100 | 4200 | 25400 | 3175 | 1075 |
| 50 | 122 | 8 | 2200 | 3800 | 23600 | 2950 | 750 |
| 50 | 123 | 8 | 2500 | 4000 | 25900 | 3238 | 738 |
| 50 | 120 | 8 | 2200 | 3200 | 22100 | 2763 | 563 |
| 50 | 124 | 8 | 2500 | 3500 | 23000 | 2875 | 375 |

Bonus Exercise (if time permits)

7. Display the sum of salary, the average of salary, and the number of employees in departments, consolidating departments 0-99 together, 100-199 together, etc.

| Depts by 100s | SUM (SALARY) | AVG (SALARY) | COUNT (*) |
|---------------|--------------|--------------|-----------|
| 0 | 612500 | 6250 | 98 |
| 100 | 71916 | 8989.5 | 8 |
| | 7000 | 7000 | 1 |

8. Display the average of all departments' average salaries. Round the result to whole currency units.

| Avg of Dept Avgs |
|------------------|
| 8153 |

9. Compare the result from the step above to the average of employee salaries. Is it the same? Why or why not?

Congratulations! You have completed this exercise.

Exercise 11.3: Using Subqueries

Connect to the HR account.

Using subqueries, write queries to display the following information:

1. Display the department id and department name for all departments that have one or more employees. Order the result by `department_id`.

```
DEPARTMENT_ID DEPARTMENT_NAME
-----
10 Administration
20 Marketing
30 Purchasing
40 Human Resources
50 Shipping
60 IT
70 Public Relations
80 Sales
90 Executive
100 Finance
110 Accounting
```

2. Display the employee id, first name, last name, and salary for all employees that have a salary greater than the average salary for all employees. Order the result by salary in descending sequence.

```
EMPLOYEE_ID FIRST_NAME LAST_NAME SALARY
-----
100 Steven King 24000
101 Neena Kochhar 17000
102 Lex De Haan 17000
145 John Russell 14000
146 Karen Partners 13500
201 Michael Hartstein 13000
205 Shelley Higgins 12008
108 Nancy Greenberg 12008
147 Alberto Errazuriz 12000
168 Lisa Ozer 11500
148 Gerald Cambrault 11000
174 Ellen Abel 11000
114 Den Raphaely 11000
162 Clara Vishney 10500
149 Eleni Zlotkey 10500
150 Peter Tucker 10000
156 Janette King 10000
204 Hermann Baer 10000
169 Harrison Bloom 10000
170 Tayler Fox 9600
163 Danielle Greene 9500
157 Patrick Sully 9500
151 David Bernstein 9500
```

Continued on next page

Exercise Manual

| | | | |
|-----|-------------|------------|------|
| 158 | Allan | McEwen | 9000 |
| 109 | Daniel | Faviet | 9000 |
| 103 | Alexander | Hunold | 9000 |
| 152 | Peter | Hall | 9000 |
| 175 | Alyssa | Hutton | 8800 |
| 176 | Jonathon | Taylor | 8600 |
| 177 | Jack | Livingston | 8400 |
| 206 | William | Gietz | 8300 |
| 110 | John | Chen | 8200 |
| 121 | Adam | Fripp | 8200 |
| 153 | Christopher | Olsen | 8000 |
| 120 | Matthew | Weiss | 8000 |
| 159 | Lindsey | Smith | 8000 |
| 122 | Payam | Kaufling | 7900 |
| 112 | Jose Manuel | Urman | 7800 |
| 111 | Ismael | Sciarra | 7700 |
| 154 | Nanette | Cambrault | 7500 |
| 160 | Louise | Doran | 7500 |
| 171 | William | Smith | 7400 |
| 172 | Elizabeth | Bates | 7300 |
| 164 | Mattea | Marvins | 7200 |
| 161 | Sarath | Sewall | 7000 |
| 155 | Oliver | Tuvault | 7000 |
| 178 | Kimberely | Grant | 7000 |
| 113 | Luis | Popp | 6900 |
| 165 | David | Lee | 6800 |
| 203 | Susan | Mavris | 6500 |
| 123 | Shanta | Vollman | 6500 |

51 rows selected.

3. Display the employee id, first name, last name, and salary for the employee that has the highest salary.

| EMPLOYEE_ID | FIRST_NAME | LAST_NAME | SALARY |
|-------------|------------|-----------|--------|
| 100 | Steven | King | 24000 |

Exercise Manual

- Display the employee id, first name, last name, salary, and commission_pct for all employees that have a salary greater than the average salary for all employees and a commission_pct greater than the average commission_pct for all employees. Order the result by last_name.

| EMPLOYEE_ID | FIRST_NAME | LAST_NAME | SALARY | COMMISSION_PCT |
|-------------|------------|-----------|--------|----------------|
| 174 | Ellen | Abel | 11000 | .3 |
| 151 | David | Bernstein | 9500 | .25 |
| 148 | Gerald | Cambrault | 11000 | .3 |
| 160 | Louise | Doran | 7500 | .3 |
| 147 | Alberto | Errazuriz | 12000 | .3 |
| 152 | Peter | Hall | 9000 | .25 |
| 175 | Alyssa | Hutton | 8800 | .25 |
| 156 | Janette | King | 10000 | .35 |
| 158 | Allan | McEwen | 9000 | .35 |
| 168 | Lisa | Ozer | 11500 | .25 |
| 146 | Karen | Partners | 13500 | .3 |
| 145 | John | Russell | 14000 | .4 |
| 161 | Sarath | Sewall | 7000 | .25 |
| 159 | Lindsey | Smith | 8000 | .3 |
| 157 | Patrick | Sully | 9500 | .35 |
| 150 | Peter | Tucker | 10000 | .3 |
| 162 | Clara | Vishney | 10500 | .25 |

17 rows selected.

Bonus Exercise (if time permits)

- Display the employee id, first name, and last name for the employee(s) that work in London. You will need to use two levels of subquery.

| EMPLOYEE_ID | FIRST_NAME | LAST_NAME |
|-------------|------------|-----------|
| 203 | Susan | Mavris |

Congratulations! You have completed this exercise.

Chapter 12: Set Operators

Exercise 12.1: Set Operators

Connect to the HR account.

1. Produce a short report showing the number of employees who earn commission and the number who do not.
 - a. Your report should look like this:

| Type | Count |
|--------------------------------------|-------|
| ----- | ----- |
| Employees who earn commission | 35 |
| Employees who do not earn commission | 72 |

- b. Use a set operator to create this report.

Congratulations! You have completed this exercise.

This page intentionally left blank.

Chapter 13: Programming with PL/SQL

Exercise 13.1: Building Anonymous Blocks

Connect to the HR account.

1. In the declaration section, declare a record `emp_rec` that uses the `employees` table as a basis.

Use %ROWTYPE.

2. In the executable section, retrieve information for employee with a last name of 'Austin' into `emp_rec`.
3. In the executable section, write a conditional structure that sets the new salary based on the employee's commission.

```
if commission_pct is undefined or zero then
    increase salary by a flat $500
if commission_pct is less than .2 then
    increase salary by $300
otherwise
    increase salary by $100
```

4. Write an `UPDATE` statement that sets the employee's salary to the value derived in the `IF` statement.
5. Select from the `employees` table for employee Austin to view the salary prior to running the block.
6. Execute the block. Again, query the `employees` table to check your results.

| LAST_NAME | SALARY |
|-----------|--------|
| ----- | ----- |
| Austin | 5300 |

7. Test your code with other employees. Use Lee and then King. Validate your results by viewing the before and after values.
8. Roll back your changes.
9. Enhance your PL/SQL block by using a `CASE` statement. Execute the block and check your results. Test with other employees as in Step 6.
10. Roll back your changes.

Bonus Exercise (if time permits)

11. Enhance the block by adding exception handling. Use the block which uses the `CASE` statement. The select statement may return no rows, one row, or many rows. Add the `EXCEPTION` section and add a handler for each of the possible errors.
12. Good practice recommends that you also add a handler for any other error that may occur. If unexpected conditions occur, raise an error, display “Contact support” and append the Oracle error message.
13. Test the program by running it for different last names. Use Austin, Smith, and Howard.
14. Roll back your changes.

Congratulations! You have completed this exercise.

Exercise 13.2: Using Cursors

Connect to the HR account.

In this exercise, you will declare and use a cursor to process the records in the `employees` table. Increase employee salary by \$5,000 for employees who were hired before Jan. 1, 2003. The cursor will accept one parameter, which limits the query to return only the required employees.

1. In the declaration section, declare a cursor that selects an employee record from the `employees` table. The cursor should accept one parameter called `in_date_hired`. Compare `in_date_hired` with the `date_hired` column in the `WHERE` clause of the `SELECT` statement. The cursor should include a `FOR UPDATE` clause to lock the selected rows.
2. Declare a record to hold the cursor results. Also, declare a numeric variable, `raise`, and initialize it to 5000. Finally, declare a date variable, `v_date`, and initialize it to Jan. 1, 2003.
3. In the executable section, open the cursor and pass `v_date` as a parameter.
4. Use a simple `LOOP...END LOOP` construct that loops through each record that the cursor returns.
5. Within the `LOOP`, use a `FETCH` statement to retrieve the row into the cursor record.
6. Add an `EXIT WHEN` statement to exit the loop when no more rows are returned by the cursor.

Warning! An exit statement must be included to avoid an endless loop.

7. Add an `UPDATE` statement, which uses the current cursor row.
8. Execute the block and check your results. Make sure to check records that should not have been updated as well!

| EMPLOYEE_ID | FIRST_NAME | LAST_NAME | HIRE_DATE | SALARY |
|-------------|------------|-----------|-------------|--------|
| 203 | Susan | Mavris | 07-JUN-2002 | 11500 |
| 204 | Hermann | Baer | 07-JUN-2002 | 15000 |
| 205 | Shelley | Higgins | 07-JUN-2002 | 17008 |
| 206 | William | Gietz | 07-JUN-2002 | 13300 |
| 102 | Lex | De Haan | 13-JAN-2001 | 22000 |
| 108 | Nancy | Greenberg | 17-AUG-2002 | 17008 |
| 109 | Daniel | Faviet | 16-AUG-2002 | 14000 |
| 114 | Den | Raphaely | 07-DEC-2002 | 16000 |

8 rows selected.

9. Roll back your changes.

Bonus Exercise (if time permits)

10. Modify the previous example to use a `FOR-LOOP` cursor instead of a regular cursor with a `LOOP`. **Set** the salary to \$11000; i.e., not a raise.

| EMPLOYEE_ID | FIRST_NAME | LAST_NAME | HIRE_DATE | SALARY |
|-------------|------------|-----------|-------------|--------|
| 203 | Susan | Mavris | 07-JUN-2002 | 11000 |
| 204 | Hermann | Baer | 07-JUN-2002 | 11000 |
| 205 | Shelley | Higgins | 07-JUN-2002 | 11000 |
| 206 | William | Gietz | 07-JUN-2002 | 11000 |
| 102 | Lex | De Haan | 13-JAN-2001 | 11000 |
| 108 | Nancy | Greenberg | 17-AUG-2002 | 11000 |
| 109 | Daniel | Faviet | 16-AUG-2002 | 11000 |
| 114 | Den | Raphaely | 07-DEC-2002 | 11000 |

8 rows selected.

11. Roll back your changes.

Congratulations! You have completed this exercise.

Chapter 14: Creating Stored Procedures, Functions, and Packages

Exercise 14.1: Stored Procedures, Functions, and Packages

Connect to the HR account.

In this exercise, you will create and execute a procedure that updates an employee if one exists. Otherwise, assume this is a new employee and insert a new employee into the table.

1. Use the `CREATE PROCEDURE` statement to define a procedure called `update_emp`. It requires six input parameters: `parm_employee_id`, `parm_last_name`, `parm_email`, `parm_hire_date`, `parm_job_id` and `parm_salary`.
2. In the executable section, update the salary in the `employees` table for the specified `employee_id`. Verify that the last name, hire date, and `job_id` match the input parameters before making the change.
3. If no rows were updated because the input employee id does not exist in the `employees` table, then insert the input data into the table as a new row.
4. Complete the procedure with an `END` statement.
5. Store the procedure in the database by executing the `CREATE PROCEDURE` statement.

If you get a warning that the procedure was created with compilation errors, use the `SHOW ERRORS` command.

6. Test the procedure by creating a simple PL/SQL block that calls the procedure and passes parameters to it. Use employees Chen and Johnston for your tests. Check your results.
7. Roll back your changes.

Bonus Exercise (if time permits)

Change the procedure from the previous exercise to perform an `INSERT` only if the department the employee is assigned to presently has a manager.

We will use a separate function to perform the test. The function is used only by the procedure. Therefore, it can be hidden by defining it as a private function in a package.

8. Start by dropping the independent procedure `update_emp` since we now want to include it into a package.
9. Write a package specification, `pack_employee` that contains a declaration of the `update_emp` procedure. Use the `CREATE PACKAGE` statement. You will need to add one more parameter to pass in a department id.
10. Use the procedure from previous section as the basis for this procedure specification. Remember to add the additional parameter for department id. Remove the procedure body, which starts with the keyword `IS` and finishes with an `END` statement.
11. Submit the package specification to the database for compilation and storage.
12. Write the package body for `pack_employee` using the `CREATE PACKAGE BODY` statement.
13. Write a function definition in the package body. The function should return the manager id for the assigned department or a null value if a manager is not assigned. The function will require a single parameter for the `department_id`.
14. The function should be defined before the procedure.
15. Use a `FUNCTION` definition statement. The `FUNCTION` consists of a `SELECT` statement and a `RETURN` statement.
16. Copy the procedure from the previous section and make the following changes:
 - a. Add an additional parameter to pass in the `department_id`.
 - b. Add an `IF-THEN` construct around the `INSERT` statement so that it inserts only if the department being assigned to has a manager assigned.
 - c. Use a `PROCEDURE` definition statement instead of the `CREATE PROCEDURE` statement within the package body.
17. The package body must be completed with an `END` statement.
18. Submit the package body to the database for compilation and storage.
19. Test the package by executing the packaged procedure from an anonymous PL/SQL block. Check your results.

Congratulations! You have completed this exercise.

Chapter 15: Testing PL/SQL

Exercise 15.1: Writing PL/SQL Tests with utPLSQL

Connect to the HR account.

In this exercise, you will create a function that checks whether a salary is appropriate for a given job. You will work TDD, writing tests using utPLSQL.

1. Start by creating a test package for your function.
 - a. Name it appropriately.
 - b. Annotate it with the `--%suite` annotation.
2. In the package specification, declare your first test specification.
 - a. Your test specification is a procedure annotated with `--%test`.
 - b. This will be a test for the normal behavior of the function.
3. Now create the package body for your test package.
 - a. Create the definition of the test specification.
 - b. It will pass in a known `job_id` and a `salary` in the right range, expecting boolean `TRUE`.
 - c. The body will not compile because the function under test does not exist.
4. Write the function.
 - a. It should accept two parameters, the `job_id`, and `salary`.
 - b. It should return `BOOLEAN`.
 - c. For now, it should just return `TRUE`.
5. Execute all tests in the current schema.
 - a. Turn on server output (`SET SERVEROUTPUT ON`).
 - b. Write an anonymous block that executes `ut.run()`. You can optionally pass in your test package name to ensure only that package runs.
 - c. One test will run, and it should pass.
6. Now write a second test in the test package.
 - a. It should be a test for a known `job_id` and a `salary` below the `min_salary`. It should expect `FALSE`.
 - b. Remember to add it to the package specification with the correct annotation.
7. Run the tests again. The new test should fail.
8. Now implement the appropriate functionality to return `FALSE` when the `salary` is below `min_salary` for the given `job_id` and `TRUE` otherwise.
 - a. For now, ignore exceptional situations, such as the `job_id` not existing, just focus on making the tests pass.

9. Run the tests and see that they pass.
10. Complete the normal behavior by testing for `FALSE` when the salary is above the `max_salary` for the given `job_id`.
 - a. Run the test, it should fail.
 - b. Write the functionality.
 - c. Re-run the test and repeat the cycle until it passes.
 - d. Refactor if necessary.
11. Now start to add some negative tests and implement the functionality. Write a single new test each time and then write the functionality to make it pass. Your function should behave like this:
 - a. If `job_id` or `salary` are `NULL`, it should raise an exception.
 - b. If `job_id` does not exist, it should return `FALSE`. Depending on how you wrote the function, you may not need to add any code to make this happen but add a test for it anyway.
 - c. If you think of any other exceptional situations, determine appropriate behavior and write a test.
12. Consider adding some additional boundary condition tests.

Congratulations! You have completed this exercise.

Exercise 15.2: Testing Updates With utPLSQL

Connect to the HR account.

In this exercise, you will create a procedure that updates a salary for an employee if the salary is appropriate for their job.

1. Start by creating a test package for your function.
 - a. Name it appropriately.
 - b. Annotate it with the `--%suite` annotation.
2. In the package specification, declare your first test specification.
 - a. Your test specification is a procedure annotated with `--%test`.
 - b. This will be a test for the normal behavior of the procedure.
3. Now create the package body for your test package.
 - a. Create the definition of the test specification.
 - b. It will pass in a known `employee_id` and a `salary` in the right range for their `job_id`, expecting a single row to be updated.
 - c. The body will not compile because the procedure under test does not exist.
4. Write the procedure.
 - a. It should accept two parameters, `employee_id` and `salary`.
 - b. For now, it should just update the row anyway.
5. Execute your new test package.
 - a. Turn on server output (`SET SERVEROUTPUT ON`).
 - b. Write an anonymous block that executes `ut.run()`. Pass in your test package name to ensure only that package runs.
 - c. One test will run, and it should pass.
6. Now write a second test in the test package.
 - a. It should be a test where the `salary` is below the `min_salary`. It should expect no rows to be updated.
 - b. Remember to add it to the package specification with the correct annotation.
7. Run the tests again. The new test should fail.
8. Now implement the appropriate functionality to ensure the `salary` is only updated when it is above the `min_salary` for the `job_id` of the employee.
 - a. If you wish, you can use the function you created in the previous exercise. Writing it as a single `UPDATE` is possible but a little harder.
 - b. For now, ignore exceptional situations, such as the `employee_id` not existing, just focus on making the tests pass.
9. Run the tests and see that they pass.

10. Complete the normal behavior by testing for no rows being updated when the salary is above the `max_salary` for the given `job_id` of the employee.
 - a. Run the test, it should fail. If you used the function in the previous steps, it may already pass. This is not a problem: the test is still valuable to protect against regression errors.
 - b. Write the functionality, if necessary.
 - c. Re-run the test and repeat the cycle until it passes.
 - d. Refactor if necessary.
 - e. Consider refactoring the tests to reduce the amount of repeated code.
11. Now start to add some negative tests and implement the functionality. Write a single new test each time, and then write the functionality to make it pass. Your procedure should behave like this:
 - a. If `employee_id` or `salary` is `NULL`, it should raise an exception.
 - b. If `employee_id` does not exist, it should simply not update any data, but it should not throw an unexpected exception (e.g., `NO_DATA_FOUND`).
 - c. If you think of any other exceptional situations, determine appropriate behavior and write a test.

Congratulations! You have completed this exercise.

Chapter 16: Creating Triggers

Exercise 16.1: Working with Triggers

Connect to the `HR` account.

In this exercise, you will write a trigger that ensures new employees are only inserted if their salary is in the right range for their job. You will use the function you created in Exercise 15.1: if you are concerned about your solution to that exercise, take the function from the solution file.

1. Check that your function has been created and that all the tests pass.
2. Create a test for inserting a new employee, expecting failure.
 - a. Note that in this case, since we are testing a trigger, the `INSERT` will occur in the test code.
 - b. Use a known unused value for the `employee_id` and `PU_CLERK` for the `job_id`. Choose a `salary` outside the right range. Set the other mandatory columns to reasonable values.
 - c. The `INSERT` should fail with an exception.
3. Run your test.
 - a. It should fail since the `INSERT` will succeed.
4. Now create an insert trigger for the `employees` table.
 - a. Use the `CREATE TRIGGER` statement.
 - b. This should be a `AFTER` trigger.
 - c. This trigger should be fired each time a row is inserted.
 - d. For now, it should just throw the exception your test is expecting.
5. Your test should now pass, but clearly, no data can be inserted into `employees`!
6. Now create a second test—this time inserting the same employee with a salary in the right range.
 - a. The test should expect the insert to succeed.
7. Run your tests. The new test should fail since the `INSERT` fails.
8. Now write the correct body of the trigger.
 - a. It should use the function to decide whether to allow the `INSERT`.
 - b. Use the `:NEW` pseudo-record to get the appropriate data values.
9. Run the tests again until they succeed.
10. When you are done, drop your trigger.

Bonus Exercise (if time permits)

11. Modify your trigger, so it also works for an `UPDATE` of `job_id` or `salary`.
 - a. Have it throw a different exception when performing an `UPDATE`.
 - b. Work TDD.
12. Drop your trigger.

Congratulations! You have completed this exercise.

Chapter 17: Data Definition Language

Exercise 17.1: Table Management

Connect to the HR account.

In this exercise, you will create a new table, a sequence, and a view. Using these, you will explore various DDL commands and their effects.

1. Create a table called `benefits`.
 - a. Use the following columns definitions:

```
benefit_id          NUMBER(3)      NOT NULL
benefit_name        VARCHAR2(25)
benefit_type        VARCHAR2(20)  DEFAULT 'HEALTH CARE'
benefit_effective_date DATE
benefit_max_allowance NUMBER(8,2)
```
 - b. Make `benefit_id` the primary key.
2. Describe the `benefits` table to verify the definition.
3. Create a sequence called `seq_benefits`. Make its starting and incremental values 1.
4. Insert a row into the `benefits` table *without* a column list.
 - a. Use the sequence for the `benefit_id`.
 - b. Make the name "401k", the type "Retirement", set the effective date to Jan. 1, 2010, and the max allowance to 250,000.
5. Insert another row into the `benefits` table *with* a column list, specifying *all* columns.
 - a. Use the sequence for the `benefit_id`.
 - b. Make the name "Medical PPO", the type "Health", set the effective date to Jan. 1, 2011, and the max allowance to 100,000.
6. Insert another row into the `benefits` table *with* a column list, specifying *all* columns.
 - a. Use the sequence for the `benefit_id`.
 - b. Set the type to the reserved word `DEFAULT`.
 - c. Make the name "Medical Ins", set the effective date to Jan. 1, 2012, and the max allowance to 125,000.
7. Display all the rows in the `benefits` table. What is the value of type for the 3rd row?

8. Insert another row into the `benefits` table with a column list. Specify all column names except for `benefit_type`.
 - a. Use the sequence for the `benefit_id`.
 - b. Make the name "No default name provided", set the effective date to Jan. 1, 2013, and the max allowance to 150,000.
9. Display all the rows in the `benefits` table. What is the value of type for the 4th row?
10. Update all benefits rows whose type value begins with "H" to the table `DEFAULT`.
11. Display all the rows in the `benefits` table. What is the value of the type columns?
12. `COMMIT` the changes.
13. Create a view called "`vw_h_b`" that contains the benefit ID, name, type, and max allowance from the `benefits` table. Only allow the rows whose value for type begins with "HEALTH".
14. Describe this view.
15. Display all the rows through the view.
16. Try to add a new, numeric, mandatory column to the `benefits` table: `max_dependents`. Why did the attempt fail?
17. Try to add the column again, this time specifying a `DEFAULT` value of 0.
18. Display the `benefits` table: what value is in the `max_dependents` column?
19. Re-run the select through the view. Does it include the new column?

Bonus Exercise (if time permits)

20. Modify the maximum size of the `benefit_name` column to be 50. Does this succeed?
 - a. Describe the `benefits` table to see the impact of the command.
21. Try to modify the maximum size of the `benefit_name` column to be 20. Why does this fail?
22. Insert into the `benefits` table by selecting all the rows from the `benefits` table.
 - a. Use the row values for all columns except for the `benefit_id`: use the sequence number for this value.
23. Display all the rows in the `benefits` table. How many are there now?

24. Issue a `ROLLBACK`.
25. Re-run the previous set insert.
 - a. Insert into the `benefits` table by selecting all the rows from the `benefits` table.
 - b. Use the row values for all columns except for the `benefit_id`: use the sequence number for this value.
26. Display all the rows in the `benefits` table. How many are there now? What are the benefit IDs? Can you explain their values?

Congratulations! You have completed this exercise.

This page intentionally left blank.

Chapter 20: Amazon DynamoDB

Exercise 20.1: Access DynamoDB Using AWS CLI

In this exercise, you will use a local version of DynamoDB and the AWS Command Line Interface to explore basic DynamoDB actions. You will first create a simple table and then add some data to it, before using various actions to query the data.

1. Use Visual Studio Code for this exercise.
 - a. Unzip the file `dynamodb.zip` into a folder on your desktop.
 - i. *Note:* running exercise from any other than the `C:\` drive will NOT work.
 - b. Open that folder in VS Code.
2. Open a terminal window and examine `start.bat`.
 - a. This file runs the local version of DynamoDB.
 - b. Run the file by entering: `.\start.bat` (*Note:* do not forget `\".\"`.)
 - c. Leave that terminal window open.
3. Open a new terminal window to use for the rest of the exercise.
4. Run the command: `aws configure`.
 - a. Do not change the Access Key or Secret Access Key, instead just press Return.
 - b. Enter a suitable region, choose `us-east-1`, `eu-west-1` or `ap-south-1`, depending on your location.
 - c. You can also press return and leave the output format unset.
5. Run the command:

```
aws dynamodb create-table --generate-cli-skeleton input
```
6. Compare the output from that command to the contents of the file `CreateTable.json`.
 - a. The output of the previous command is a pro forma for the `CreateTable` action.
 - b. The file `CreateTable.json` creates a table called `Music` with a two-part key (`artist` and `songTitle`).

7. Run the following command and check the output:

```
aws dynamodb create-table \  
  --endpoint-url http://localhost:8000 \  
  --cli-input-json file://CreateTable.json
```

- Replace the backslash (\) with the appropriate line continuation character for your terminal or type the command on one line.
- Make sure not to forget the `--endpoint-url` parameter or the `file://` protocol on the `--cli-input-json` parameter.
- Normally, you would need to wait for the table to become active, but since we are using the local version of DynamoDB, it becomes active immediately.

8. Run these commands to see your new table:

```
aws dynamodb list-tables \  
  --endpoint-url http://localhost:8000  
  
aws dynamodb describe-table \  
  --endpoint-url http://localhost:8000 \  
  --table-name Music
```

9. Run the command:

```
aws dynamodb put-item --generate-cli-skeleton input
```

- Again, the output is a pro forma for the `PutItem` action.
- Note that this pro forma contains a number of legacy items. Consult the documentation to see which items those are:
https://docs.aws.amazon.com/amazondynamodb/latest/APIReference/API_PutItem.html

10. Add some data to your table:

```
aws dynamodb put-item \  
  --endpoint-url http://localhost:8000 \  
  --cli-input-json file://PutItem1.json  
  
aws dynamodb put-item \  
  --endpoint-url http://localhost:8000 \  
  --cli-input-json file://PutItem2.json
```

- Note that the two files do not set the same attributes on both items.

11. Try different ways of retrieving data (review each file before you run it):

```
aws dynamodb get-item \  
  --endpoint-url http://localhost:8000 \  
  --cli-input-json file://GetItem.json  
  
aws dynamodb query \  
  --endpoint-url http://localhost:8000 \  
  --cli-input-json file://Query.json  
  
aws dynamodb scan \  
  --endpoint-url http://localhost:8000 \  
  --cli-input-json file://Scan.json
```

Bonus Exercise (if time permits)

12. Try other variations of these commands; for example, you can also insert an item using this version of the `PutItem` action:

```
aws dynamodb put-item \  
  --endpoint-url http://localhost:8000  
  --table-name Music \  
  --item file://item.json \  
  --return-values ALL_OLD
```

- a. It won't return any data unless you run it twice because the contents of `item.json` are a new item.
 - b. You could also specify the item at the command line in either JSON or shorthand format.
13. Using your own data, add new items to the `Music` table and write `GetItem`, `Query`, and `Scan` commands to retrieve them.

Congratulations! You have completed this exercise.

Exercise 20.2: Java Document API

In this exercise, you will access DynamoDB from Java code and investigate how to achieve the same basic tasks using the Java Document API.

1. Make sure the local version of DynamoDB is running, as described at the start of the previous exercise.
2. Use Eclipse for the rest of the exercise.
3. Open the `DynamoDB` project in the Relational Databases workspace.
4. Open `DynamoDbDocumentDao` and `DynamoDbDocumentDaoTest`.
 - a. Review the tests and corresponding code.
 - b. Compare them to the files used for the previous exercise.
5. Run the tests. They should all pass.
6. Take a copy of `music.json`.
 - a. Do not edit the existing file since it is used by a number of test classes.
 - b. Change the contents of your new JSON file to reflect some of your own musical choices.
 - c. Edit `DynamoDbDocumentDaoTest` to refer to the new file.
 - d. Make the tests pass with your new data.

Bonus Exercise (if time permits)

7. Open `DynamoDbLowLevelDao` and `DynamoDbDaoLowLevelTest`. These use the low-level API.
 - a. Review the code and compare them to the Document versions.
 - b. Especially look at features that you know to be different, such as waiting for asynchronous operations, converting to and from Java classes, and iterating over windowed collections.

Congratulations! You have completed this exercise.

Exercise 20.3: Java Object Mapper

In this exercise, you will continue accessing DynamoDB using Java, but this time you will access data using the Java Object Mapper. You will see how to achieve tasks such as inserting and querying data using this higher-level interface.

1. Make sure the local version of DynamoDB is running, as described at the start of the previous exercise.
2. In the `DynamoDB` project, open `DynamoDbMapperDao` and `DynamoDbMapperDaoTest`.
 - a. Review the tests and corresponding code.
 - b. Compare them to the files used for the previous exercise.
3. Run the tests. They should all pass.
4. Take a copy of `music.json`.
 - a. Do not edit the existing file since it is used by a number of test classes.
 - b. Change the contents of your new JSON file to reflect some of your own musical choices.
 - c. Edit `DynamoDbMapperDaoTest` to refer to the new file.
 - d. Make the tests pass with your new data.

Congratulations! You have completed this exercise.